

Verifiable C

*Applying the Verified Software Toolchain
to C programs*

*Version 1.5+
December 29, 2015*

Andrew W. Appel

with Josiah Dodds and Qinxiang Cao

Contents

Verifiable C	i
Contents	ii
1 Overview	5
2 Installation	7
3 Clightgen and ASTs	8
4 Use the IDE	10
5 Functional spec, API spec	11
6 Proof of the sumarray program	16
7 start_function	17
8 forward	19
9 While loops	21
10 Entailments	23
11 Array subscripts	26
12 Splitting sublists	28
13 Returning from a function	30
14 Global variables and main()	31
15 Tying all the functions together	33
16 Separation logic: EX, *, emp, !!	34
17 PROP() LOCAL() SEP()	35
18 EX, Intros, Exists	36
19 Integers: nat, Z, int	38
20 Values: Vint, Vptr	40
21 C types	41
22 CompSpecs	42
23 reptype	43
24 Uninitialized data, default_val	44
25 data_at	45
26 reptype', repinj	46
27 field_at	48
28 Localdefs: temp, lvar, gvar	49
29 go_lower	50

30 saturate_local	51
31 field_compatible, field_address	52
32 value_fits	54
33 Cancel	55
34 entailer!	56
35 Shares	57
36 Pointer comparisons	59
37 Proving larg(ish) programs	60
38 Separate compilation, semax_ext	63
39 Appendix: catalog of major tactics/commands	64
40 BOUNDARY	68
41 Getting started	69
42 Differences from PLCC	71
43 Memory predicates	72
44 Separation Logic	73
45 Mapsto and func_ptr	75
46 CompCert C	76
47 Verifiable C programming	77
48 32-bit Integers	78
49 C expression syntax	81
50 C operators	82
51 C expression evaluation	85
52 C type checking	87
53 Lifted separation logic	88
54 Canonical forms	90
55 Supercanonical forms	91
56 Go_lower	92
57 Welltypedness of variables	93
58 Normalize	95
59 Entailer	100
60 The Hoare triple	101
61 Later	102

62 Specifying a function	103
63 Specifying all functions	105
64 Proving a function	106
65 Manipulating preconditions	109
66 The Frame rule	111
67 Pointer comparisons	112
68 Structured data	113
69 For loops	117
70 Nested Loads	119
71 JUNK	123

1 Overview

Verifiable C is a language and program logic for reasoning about the functional correctness of C programs. The *language* is a subset of CompCert C light; it is a dialect of C in which side-effects and loads have been factored out of expressions. The *program logic* is a higher-order separation logic, a kind of Hoare logic with better support for reasoning about pointer data structures, function pointers, and data abstraction.

Verifiable C is *foundationally sound*. That is, it is proved (with a machine-checked proof in the Coq proof assistant) that,

Whatever observable property about a C program you prove using the Verifiable C program logic, that property will actually hold on the assembly-language program that comes out of the C compiler.

This soundness proof comes in two parts: The program logic is proved sound with respect to the semantics of CompCert C, by a team of researchers primarily at Princeton University; and the C compiler is proved correct with respect to those same semantics, by a team of researchers primarily at INRIA. There are other proofs above (regarding proof automation tools for the use of the program logic) and below (regarding the execution of assembly language on a concurrent machine). This chain of proofs from top to bottom, connected in Coq at specification interfaces, is called the *Verified Software Toolchain*.



To use Verifiable C, one must have had some experience using Coq, and some familiarity with the basic principles of Hoare logic. These can be obtained by studying Pierce’s *Software Foundations* interactive textbook, and doing the exercises all the way to chapter “Hoare2.”

It is also useful to read the brief introductions to Hoare Logic and Separation Logic, covered in Appel’s *Program Logics for Certified Compilers*, [Chapters 2 and 3](#).

2 Installation

The Verified Software Toolchain runs on Linux, Mac, or Windows. You will need to install:

1. Coq 8.4pl6, from coq.inria.fr. Follow the standard installation instructions.
2. CompCert 2.5 (or 2.5.1 if available), from compcert.inria.fr. You will want to build the *clightgen* tool, using these commands: `./configure ia32-linux; make clightgen`. You might replace `ia32-linux` with `ia32-macosx` or `ia32-cygwin`.
3. VST 1.6 (if available), or else a recent stable trunk version such as [f42918bceb](#). Follow the instructions in the file `BUILD_ORGANIZATION`.

WORKFLOW. Within `vst`, the `progs` directory contains some sample C programs with their verifications. The workflow is:

- Write a C program *F.c*.
- Run `clightgen F.c` to translate it into a Coq file *F.v*.
- Write a verification of *F.v* in a file such as `verif_F.v`. That latter file must import both *F.v* and the VST *Floyd*¹ program verification system, `floyd.proofauto`.

LOAD PATHS. Interactive development environments (CoqIDE or Proof General) will need their load paths properly initialized through command-line arguments. Running `make` in `vst` creates a file `.loadpath` with the right arguments. You can then do (for example),

```
coqide `cat .loadpath` progs/verif_reverse.v
```

See the heading USING PROOF GENERAL AND COQIDE in the file `BUILD_ORGANIZATION` for more information.

¹Named after Robert W. Floyd (1936–2001), a pioneer in program verification.

3 *Clightgen and ASTs*

We will introduce Verifiable C by explaining the proof of a simple C program: adding up the elements of an array.

```
#include <stddef.h>
```

```
int sumarray(int a[], int n) {
  int i,s,x;
  i=0;
  s=0;
  while (i<n) {
    x=a[i];
    s+=x;
    i++;
  }
  return s;
}
```

```
int four[4] = {1,2,3,4};
```

```
int main(void) {
  int s;
  s = sumarray(four,4);
  return s;
}
```

You can examine this program in `VST/progs/sumarray.c`. Then look at `progs/sumarray.v` to find the output of CompCert's *clightgen* utility: it is the abstract syntax tree (AST) of the C program, expressed in Coq. In `sumarray.v` there are definitions such as,

...

Definition `_main` : ident := 54%positive.

...

Definition `_s` : ident := 50%positive.

...

Definition $f_sumarray := \{ |$
 $fn_return := tint; \dots$
 $fn_params := ((_a, (tptr\ tint)) :: (_n, tint) :: nil);$
 $fn_temps := ((_i, tint) :: (_s, tint) :: (_x, tint) :: nil);$
 $fn_body :=$
 $(Ssequence$
 $(Sset\ _i\ (Econst_int\ (Int.repr\ 0)\ tint))$
 $(Ssequence$
 $(Sset\ _s\ (Econst_int\ (Int.repr\ 0)\ tint))$
 $(Ssequence\ \dots$
 $)))$
 $| \}$.

...

Definition $prog : Clight.program := \{ | \dots | \}$

In general it's never necessary to read the AST file such as `sumarray.v`. But it's useful to know what kind of thing is in there. C-language identifiers such as `main` and `s` are represented in ASTs as positive numbers; the definitions `_main` and `_s` are abbreviations for these. The AST for `sumarray` is in the function-definition `f_sumarray`.

There you can see that `sumarray`'s return type is `int`. To represent the syntax of C type-expressions, CompCert defines,

Inductive $type : Type :=$
 $| Tvoid: type$
 $| Tint: intsize \rightarrow signedness \rightarrow attr \rightarrow type$
 $| Tpointer: type \rightarrow attr \rightarrow type$
 $| Tstruct: ident \rightarrow attr \rightarrow type$
 $| \dots$

and we abbreviate $tint := Tint\ l32\ Signed\ noattr$.

4 *Use the IDE*

[Chapter 5](#) through [Chapter 15](#) are meant to be read while you have the file `progs/verif_sumarray.v` open in a window of your interactive development environment for Coq. You can use Proof General, CoqIDE, or any other IDE that supports Coq.

Reading these chapters will be much less informative if you cannot see the proof state as each chapter discusses it.

Before starting the IDE, read about load paths, at the heading `USING PROOF GENERAL AND COQIDE` in the file `VST/BUILD_ORGANIZATION`.

5 *Functional spec, API spec*

A program without a specification cannot be incorrect, it can only be surprising.
(Paraphrase of J. J. Horning, 1982)

The file `progs/verif_sumarray.v` contains the specification of `sumarray.c`, and the proof of correctness of the C program with respect to that specification. For larger programs, one would typically break this down into three or more files:

1. Functional specification
2. API specification
3. Function-body correctness proofs, one per file.

To prove correctness of `sumarray.c`, we start by writing a *functional spec* of adding-up-a-sequence, then an *API spec* of adding-up-an-array-in-C.

FUNCTIONAL SPEC. A *mathematical model* of this program is the sum of a sequence of integers: $\sum_{i=0}^{n-1} x_i$. It's conventional in Coq to use `list` to represent a sequence; we can represent the sum with a list-fold:

Definition `sum_Z` : `list Z` → `Z` := `fold_right Z.add 0`.

A functional spec contains not only definitions; it's also useful to include theorems about this mathematical domain:

Lemma `sum_Z_app`: $\forall a\ b, \text{sum_Z } (a++b) = \text{sum_Z } a + \text{sum_Z } b$.

Proof.

intros. induction a; simpl; omega.

Qed.

The data types used in a functional spec can be any kind of mathematics at all, as long as we have a way to relate them to the integers, tuples, and sequences used in a C program. But the mathematical integers `Z` and the 32-bit modular integers `Int.int` are often relevant. Notice that this functional spec does not depend on `sumarray.v` or even on anything in the

Verifiable C libraries. This is typical, and desirable: the functional spec is about mathematics, not about C programming.

THE APPLICATION PROGRAMMER INTERFACE of a C program is expressed in its header file: function prototypes and data-structure definitions that explain how to call upon the modules' functionality. In *Verifiable C*, an *API specification* is written as a series of *function specifications* (funspecs) corresponding to the function prototypes.

We start `verif_sumarray.v` with some standard boilerplate:

Require Import floyd.proofauto.

Require Import progs.sumarray.

Instance CompSpecs : compspecs. make_compspecs prog. **Defined.**

Definition Vprog : varspecs. mk_varspecs prog. **Defined.**

The first line imports Verifiable C and its *Floyd* proof-automation library. The second line imports the AST of the program to be proved. Lines 3 and 4 are identical in any verification: see [Chapter 22](#) and `??`.

After the boilerplate (and the functional spec), we have the function specifications for each function in the API spec:

Definition sumarray_spec :=

DECLARE _sumarray

WITH a: val, sh : share, contents : list Z, size: Z

PRE [_a OF (tptr tint), _n OF tint]

PROP(readable_share sh;

0 ≤ size ≤ Int.max_signed;

Forall (fun x ⇒ Int.min_signed ≤ x ≤ Int.max_signed) contents)

LOCAL(temp _a a; temp _n (Vint (Int.repr size)))

SEP(data_at sh (tarray tint size) (map Vint (map Int.repr contents)) a)

POST [tint]

PROP()

LOCAL(temp ret_temp (Vint (Int.repr (sum_Z contents))))

SEP(data_at sh (tarray tint size) (map Vint (map Int.repr contents)) a).

The funspec begins, **Definition** $f_spec := \text{DECLARE } id_f \dots$ where f is the name of the C function.

A function is specified by its *precondition* and its *postcondition*. The WITH clause quantifies over Coq values that may appear in both the precondition and the postcondition. The precondition is parameterized by the C-language function parameters, and the postcondition is parameterized by a identifier `ret_temp`, which is short for, “the temporary variable holding the return value.” But really, the Coq variable `_a` does not have type (pointer-to-int); it has type `ident` (see [page 8](#)).

An assertion in Verifiable C’s *separation logic* can be written at either of two levels: The *lifted level*, implicitly quantifying over all local-variable states; or the *base level*, at a particular local-variable state. Program assertions are written at the lifted level, for which the notation is `PROP(...) LOCAL(...) SEP(...)`.

In an assertion `PROP(\vec{P}) LOCAL(\vec{Q}) SEP(\vec{R})`, the propositions in the sequence \vec{P} are all of Coq type `Prop`. They describe things that are forever true, independent of program state. Of course, in the function precondition above, the statement `0 ≤ size ≤ Int.max_signed` is “forever” true *just within the scope of the quantification of the variable size*; it is bound by WITH and spans the PRE and POST assertions.

The LOCAL propositions \vec{Q} are *variable bindings* of type `localdef`. Here, the function-parameters a and n are treated as nonaddressable local variables, or “temp” variables. The `localdef (temp _a a)` says that (in this program state) the contents of C local variable `_a` is the Coq value a . In general, the contents of a C scalar variable is always a `val`; this type is defined by CompCert as,

Inductive `val`: `Type := Vundef: val | Vint: int → val | Vlong: int64 → val`
`| Vfloat: float → val | Vsingle: float32 → val | Vptr: block → int → val.`

The SEP conjuncts \vec{R} are *spatial assertions* in separation logic. In this

case, there's just one, a `data_at` assertion saying that at address `a` in memory, there is a data structure of type *array[size] of integers*, with access-permission `sh`, and the contents of that array is the sequence map `Vint` contents.

THE POSTCONDITION is introduced by `POST [tint]`, indicating that this function returns a value of type `int`. There are no `PROP` statements in the postcondition, because no forever-true facts exist in the world that weren't already true on entry to the function. (This is typical!) The `LOCAL` *must not mention* the function parameters, because they are destroyed on function exit; it will only mention the return-temporary `ret_temp`. The `SEP` clause mentions all the spatial resources from the precondition, minus ones that have been freed (deallocated), plus ones that have been malloc'd (allocated).

So, overall, the specification for `sumarray` is this: “At any call to `sumarray`, there exist values *a, sh, contents, size* such that *sh* gives at least read-permission; *size* is representable as a nonnegative 32-bit signed integer; function-parameter `_a` contains value *a* and `_n` contains the 32-bit representation of *size*; and there's an array in memory at address *a* with permission *sh* containing *contents*. The function returns a value equal to `sum_int(contents)`, and leaves the array unaltered.”

INTEGER OVERFLOW. The C language specification says that a C compiler *may* treat signed integer overflow by wrapping around mod 2^n , where n is the word size (e.g., 32). In practice, almost all C compilers (including CompCert) do this wraparound, and it is part of the CompCert C light operational semantics. See [Chapter 19](#). The function `Int.repr`: $\mathbb{Z} \rightarrow \text{int}$ truncates mathematical integers into 32-bit integers by taking the (sign-extended) low-order 32 bits. `Int.signed`: $\text{int} \rightarrow \mathbb{Z}$ injects back into the signed integers.

The postcondition guarantees that the value return is `Int.repr (sum_Z contents)`. But what if $\sum s \geq 2^{31}$, so the sum doesn't fit in a 32-bit signed integer? Then `Int.signed(Int.repr (sum_Z contents)) \neq (sum_Z contents)`. In gen-

eral, for a claim about $\text{Int.repr}(x)$ to be *useful*, one also needs a claim that $0 \leq x \leq \text{Int.max_unsigned}$ or $\text{Int.min_signed} \leq x \leq \text{Int.max_signed}$. The caller of this function will probably need to prove $\text{Int.min_signed} \leq \text{sum_Z contents} \leq \text{Int.max_signed}$ in order to make much use of the post-condition.

What if s is the sequence $[\text{Int.max_signed}; 5; 1 - \text{Int.max_signed}]$? Then $\sum s = 6$. Does the program really work? Answer: Yes, by the miracle of modular arithmetic.

6 *Proof of the sumarray program*

To prove correctness of a whole program,

1. Collect the function-API specs together into Gprog: list funspec.
2. Prove that each function satisfies its own API spec (with a `semax_body` proof).
3. Tie everything together with a `semax_func` proof.

In `progs/verif_sumarray.v`, the first step is easy:

Definition `Gprog : funspecs := sumarray_spec :: main_spec::nil`.

The function specs, built using `DECLARE`, are listed in the same order the functions appear in the program (in particular, the same order they appear in `prog.(prog_defs)`, in `sumarray.v`).

In addition to `Gprog`, the API spec contains `Vprog`, the list of global-variable type-specs. This is computed automatically by the `mk_varsspecs` tactic, as shown at the beginning of `verif_sumarray.v`.

Each C function can call any of the other C functions in the API, so each `semax_body` proof is a client of the entire API spec, that is, `Vprog` and `Gprog`. You can see that in the statement of the `semax_body` lemma for the `_sumarray` function:

Lemma `body_sumarray: semax_body Vprog Gprog f_sumarray sumarray_spec`.

Here, `f_sumarray` is the actual function body (AST of the C code) as parsed by `clightgen`; you can read it in `sumarray.v`. You can read `body_sumarray` as saying, *In the context of `Vprog` and `Gprog`, the function body `f_sumarray` satisfies its specification `sumarray_spec`*. We need the context in case the `sumarray` function refers to a global variable (`Vprog` provides the variable's type) or calls a global function (`Gprog` provides the function's API spec).

7 start_function

The predicate `semax_body` states the Hoare triple of the function body, $\Delta \vdash \{Pre\}c\{Post\}$. *Pre* and *Post* are taken from the funspec for *f*, *c* is the body of *F*, and the type-context Δ is calculated from the global type-context overlaid with the parameter- and local-types of the function.

To prove this, we begin with the tactic `start_function`, which takes care of some simple bookkeeping and expresses the Hoare triple to be proved.

Lemma `body_sumarray`: `semax_body Vprog Gprog f_sumarray sumarray_spec`.
Proof.

`start_function`.

The proof goal now looks like this:

```

Espec : OracleKind
a : val
sh : share
contents : list Z
size : Z
Delta_specs := abbreviate : PTree.t funspec
Delta := abbreviate : tycontext
SH : readable_share sh
H : 0 ≤ size ≤ Int.max_signed
H0 : Forall (fun x : Z ⇒ Int.min_signed ≤ x ≤ Int.max_signed) contents
POSTCONDITION := abbreviate : ret_assert
MORE_COMMANDS := abbreviate : statement
----- (1/1)
semax Delta
  (PROP ()
    LOCAL(temp _a a; temp _n (Vint (Int.repr size)))
    SEP(data_at sh (tarray tint size) (map Vint (map Int.repr contents)) a))
  (Ssequence (Sset _i (Econst.int (Int.repr 0) tint)) MORE_COMMANDS)
  POSTCONDITION

```

First we have *Espec*, which you can ignore for now (it characterizes the outside world, but `sumarray.c` does not do any I/O). Then `a,sh,contents,size` are exactly the variables of the `WITH` clause of `sumarray_spec`.

The two abbreviations `Delta_spec`, `Delta` are the type-context in which Floyd's proof tactics will look up information about the types of the program's variables and functions. The hypotheses `SH,H,H0` are exactly the `PROP` clause of `sumarray_spec`'s precondition. The `POSTCONDITION` is exactly the `POST` part of `sumarray_spec`.

To see the contents of an abbreviation, you can either (1) set your IDE to show implicit arguments, or (2) (e.g.) unfold abbreviate in `POSTCONDITION`.

Below the line we have one proof goal: the Hoare triple of the function body. It's written, $\text{semax} \Delta P \ c \ Q$, where P is the precondition, c is the command, and Q is the postcondition. Because we do *forward* Hoare-logic proof, we won't care about the postcondition until we get to the end of c , so here we hide it away in an abbreviation. Here, the command c is a long sequence starting with `i=0;...more`, and we hide the *more* in an abbreviation `MORE_COMMMANDS`.

The precondition of this `semax` has `LOCAL` and `SEP` parts taken directly from the `funspec` (the `PROP` clauses have been moved above the line). The statement `(Sset _i (Econst_int (Int.repr 0) tint))` is the AST generated by `clightgen` from the C statement `i=0;.`

8 forward

We do Hoare logic proof by forward symbolic execution. On [page 17](#) we show the proof goal at the beginning of the `sumarray` function body. In a forward Hoare logic proof of $\{P\} i = 0; \text{more} \{R\}$ we might first apply the sequence rule,

$$\frac{\{P\} i = 0 \{Q\} \quad \{Q\} \text{more} \{R\}}{\{P\} i = 0; \text{more} \{R\}}$$

assuming we could derive some appropriate assertion Q .

For many kinds of statements (assignments, return, break, continue) this is done automatically by the forward tactic. When we execute forward here, the resulting proof goal is,

Espec, a, sh, contents, size, Delta_spec, SH, H, H0 *as before*

Delta := abbreviate : tycontext

POSTCONDITION := abbreviate : ret_assert

MORE_COMMANDS := abbreviate : statement

----- (1/1)

semax Delta

(PROP ())

LOCAL(temp _i (Vint (Int.repr 0)); temp _a a;

temp _n (Vint (Int.repr size)))

SEP(data_at sh (tarray tint size) (map Vint (map Int.repr contents)) a))

(Ssequence (Sset _s (Econst_int (Int.repr 0) tint)) MORE_COMMANDS)

POSTCONDITION

Notice that the precondition of this `semax` is really the *postcondition* of the `i=0;` statement; it is the precondition of the *next* statement, `s=0;`. It's much like the precondition of `i=0;` what has changed?

- The `LOCAL` part contains `temp _i (Vint (Int.repr 0))` in addition to what it had before; this says that the local variable `i` contains integer value zero.

- the command is now `s=0;more`, where `MORE_COMMANDS` no longer contains `s=0;`.
- Delta has changed; it now records the information that i is initialized.

Another forward goes through `s=0;` to yield a proof goal with a `LOCAL` binding for the `_s` variable.

9 While loops

To prove a *while* loop by forward symbolic execution, you use the tactic `forward_while`, and you must supply a loop invariant. Take the example of the `forward_while` in `progs/verif_sumarray.v`. The proof goal is,

```

Espec, Delta_specs, Delta
a : val, sh : share, contents : list Z, size : Z
SH : readable_share sh
H : 0 ≤ size ≤ Int.max_signed
H0 : Forall (fun x : Z ⇒ Int.min_signed ≤ x ≤ Int.max_signed) contents
POSTCONDITION := abbreviate : ret_assert
MORE_COMMANDS, LOOP_BODY := abbreviate : statement
-----(1/1)

```

```

semax Delta
(PROP ()
  LOCAL(temp _s (Vint (Int.repr 0)); temp _i (Vint (Int.repr 0));
    temp _a a; temp _n (Vint (Int.repr size)))
  SEP(data_at sh (tarray tint size) (map Vint (map Int.repr contents)) a))
(Ssequence
  (Swhile (Ebinop Olt (Etempvar _i tint) (Etempvar _n tint) tint)
    LOOP_BODY)
  MORE_COMMANDS)
POSTCONDITION

```

A loop invariant is an assertion, almost always in the form of an existential `EX...PROP()LOCAL()SEP()`. Each iteration of the loop has a state characterized by a different value of some iteration variable(s), the `EX` binds that value. For example, the invariant for this loop is,

Definition `sumarray_Inv a0 sh contents size :=`

```

EX i: Z,
  PROP(0 ≤ i ≤ size)
  LOCAL(temp _a a0; temp _i (Vint (Int.repr i)); temp _n (Vint (Int.repr size));
    temp _s (Vint (Int.repr (sum_Z (sublist 0 i contents)))))
  SEP(data_at sh (tarray tint size) (map Vint (map Int.repr contents)) a0).

```

The existential binds i , the iteration-dependent value of the local variable named $_i$. In general, there may be any number of EX quantifiers.

1. the precondition (of the whole loop) implies the loop invariant;
2. the loop-condition expression type-checks (i.e., guarantees to evaluate successfully);
3. the postcondition of the loop body implies the loop invariant;
4. the loop invariant (and *not* loop condition) is a good precondition for the proof of the MORE_COMMANDS after the loop.

Let's take a look at that first subgoal:

(above-the-line hypotheses elided) 1/4

ENTAIL Delta,
 PROP()
 LOCAL(temp $_s$ (Vint (Int.repr 0)); temp $_i$ (Vint (Int.repr 0));
 temp $_a$ a; temp $_n$ (Vint (Int.repr size)))
 SEP(data_at sh (tarray tint size) (map Vint (map Int.repr contents)) a)
 \vdash EX $i : \mathbb{Z}$,
 PROP($0 \leq i \leq \text{size}$)
 LOCAL(temp $_a$ a; temp $_i$ (Vint (Int.repr i));
 temp $_n$ (Vint (Int.repr size));
 temp $_s$ (Vint (Int.repr (sum_Z (sublist 0 i contents)))))
 SEP(data_at sh (tarray tint size) (map Vint (map Int.repr contents)) a)

This is an *entailment* goal; [Chapter 10](#) shows how to prove such goals.

10 Entailments

An *entailment* in separation logic, $P \vdash Q$, says that any state satisfying P must also satisfy Q . What’s in a state? Local-variable environment, heap (addressable memory), even the state of the outside world. VST’s type `mpred`, *memory predicate*, can be thought of as $\text{mem} \rightarrow \text{Prop}$ (but is not quite the same, for quite technical semantic reasons). That is, an `mpred` is a test on the heap only, and cannot “see” the local variables (tempvars) of the C program.

Type `environ` is a local/global variable environment, mapping identifiers (`ident`) to the values of globals, addressable locals, and tempvars (nonaddressable locals). A *lifted predicate* of type $\text{type environ} \rightarrow \text{mpred}$ can “see” both the heap and the local/global variables. The Pre/Post arguments of Hoare triples (`semax` Δ Pre `c` Post) are lifted predicates.

At present, Verifiable C has a notion of external-world state, in the `Espec`: `OracleKind`, but it is not well developed; enhancements will be needed for reasoning about input/output.

Our language for lifted predicates uses $\text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(\vec{R})$, where \vec{R} is a list of `mpreds`. Our language for `mpreds` uses primitives such as `data_at` and `emp`, along with connectives such as the $*$ and $\neg*$ of separation logic. In both languages there is an `EX` operator for existential quantification.

Separation logic’s rule of consequence is shown here

$$\frac{P \vdash P' \quad \{P'\} c \{Q'\} \quad Q' \vdash Q}{\{P'\} c \{Q'\}} \quad \frac{\Delta, P \vdash P' \quad \text{semax } \Delta P' c Q' \quad \Delta, Q' \vdash Q}{\text{semax } \Delta P c Q}$$

at left in traditional notation, and at right as in Verifiable C. The type-context Δ constrains values of locals and globals. Using this axiom, called `semax_pre_post` on a proof goal `semax` $\Delta P c Q$ yields three subgoals: another `semax` and two (lifted) entailments, $\Delta, P \vdash P'$ and $\Delta, Q' \vdash Q'$.

The standard form of a lifted entailment is $\text{ENTAIL } \Delta, \text{PQR} \vdash \text{PQR}'$, where PQR and PQR' are typically in the form $\text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(\vec{R})$, perhaps with some EX quantifiers in the front. The turnstile \vdash is written in Coq as $|--$.

Let's consider the entailment arising from `forward_while` in the `progs/verif_sumarray` example:

$$\frac{\begin{array}{l} \text{H : } 0 \leq \text{size} \leq \text{Int.max_signed} \\ \text{(other above-the-line hypotheses elided)} \end{array}}{\text{ENTAIL Delta,}}_{1/4}$$

$$\begin{array}{l} \text{PROP()} \\ \text{LOCAL(temp_s (Vint (Int.repr 0)); temp_i (Vint (Int.repr 0));} \\ \quad \text{temp_a a; temp_n (Vint (Int.repr size)))} \\ \text{SEP(data_at sh (tarray tint size) (map Vint (map Int.repr contents)) a)} \\ \vdash \text{EX } i : \mathbb{Z}, \\ \quad \text{PROP}(0 \leq i \leq \text{size}) \\ \quad \text{LOCAL(temp_a a; temp_i (Vint (Int.repr i));} \\ \quad \quad \text{temp_n (Vint (Int.repr size));} \\ \quad \quad \text{temp_s (Vint (Int.repr (sum_Z (sublist 0 i contents)))))} \\ \quad \text{SEP(data_at sh (tarray tint size) (map Vint (map Int.repr contents)) a)} \end{array}$$

We instantiate the existential with the only value that works here, zero: **Exists 0**. [Chapter 18](#) explains how to handle existentials with **Intros** and **Exists**.

Now we use the `entailer!` tactic to solve as much of this goal as possible (see [Chapter 59](#)). In this case, the goal solves entirely automatically. In particular, $0 \leq i \leq \text{size}$ solves by `omega`; `sublist 0 0 contents` rewrites to `nil`; and `sum_Z nil` simplifies to 0.

THE SECOND SUBGOAL of `forward_while` in `progs/verif_sumarray.v` is a *type-checking entailment*, of the form $\text{ENTAIL } \Delta, \text{PQR} \vdash \text{tc_expr } \Delta \ e$ where e is (the abstract syntax of) a C expression; in the particular case of a *while* loop, e is the negation of the loop-test expression. The

entailment guarantees that e executes without crashing: all the variables it references exist, and are initialized; and it doesn't divide by zero, et cetera.

In this case, the entailment concerns the expression $\neg(i < n)$,

ENTAIL Delta, PROP(...) LOCAL(...) SEP(...)

⊢ tc_expr Delta

(Eunop Onotbool (Ebinop Olt (Etempvar _i tint) (Etempvar _n tint) tint)
tint)

This solves completely via the `entailer!` tactic. To see why that is, instead of doing `entailer!`, do `unfold tc_expr; simpl`. You'll see that the right-hand side of the entailment simplifies down to `!!True`. That's because the typechecker is *calculational*, as Chapter 25 of *Program Logics for Certified Compilers* explains.

11 Array subscripts

THE THIRD SUBGOAL of `forward_while` in `progs/verif_sumarray.v` is the *body* of the while loop: `{x=a[i]; s+=x; i++;}`.

This can be handled by three forward commands, but the first one of these leaves a subgoal—proving that the subscript i is in range. Let's examine the proof goal:

SH : readable_share sh

H : $0 \leq \text{size} \leq \text{Int.max_signed}$

H0 : Forall (fun $x : Z \Rightarrow \text{Int.min_signed} \leq x \leq \text{Int.max_signed}$) contents

$i : Z$

HRE : $i < \text{size}$

H1 : $0 \leq i \leq \text{size}$

----- (1/1)

semex Delta

(PROP ())

LOCAL(temp _a a ; temp _i (Vint (Int.repr i));

temp _n (Vint (Int.repr size));

temp _s (Vint (Int.repr (sum_Z (sublist 0 i contents))))))

SEP(data_at sh (tarray tint size) (map Vint (map Int.repr contents)) a))

(Ssequence

(Sset _x

(Ederef

(Ebinop Oadd (Etempvar _a (tptr tint)) (Etempvar _i tint)

(tptr tint)) tint)) MORE_COMMANDS) POSTCONDITION

The Coq variable i was introduced automatically by `forward_while` from the existential variable, the EX $i:Z$ of the loop invariant.

The command `x=a[i];` is a *load* from data-structure a . For this to succeed, there must be a `data_at` (or `field_at`) assertion about a in the SEP clauses of the precondition; the permission share in that `data_at` must grant read access; and the subscript must be in range. Indeed, the `data_at` is there,

and the share is taken care of automatically by the hypothesis SH above the line.

So, forward succeeds; but it leaves an array-bounds subgoal:

```
ENTAIL Delta, PROP(...) LOCAL(...) SEP(...)
├ tc_expr Delta (Etempvar _a (tptr tint)) &&
  local `(tc_val tint (Znth i (map Vint (map Int.repr contents)) Vundef)) &&
    (tc_expr Delta (Etempvar _i tint) && TT)
```

The two `tc_expr` conjuncts are trivial (they are $\beta\eta$ -equal to `TT`) but the middle conjunct is nontrivial. To clean things up, we run `entailer!`, which leaves this subgoal:

```
HRE : i < Zlength (map Vint (map Int.repr contents))
H1 : 0 ≤ i ≤ Zlength (map Vint (map Int.repr contents))
  (other above-the-line hypotheses elided)
────────────────────────────────────────────────────────────────────────────────
is_int I32 Signed (Znth i (map Vint (map Int.repr contents)) Vundef)
```

For the load to succeed, the i element of `(map Vint (map Int.repr contents))` must actually be an integer, not an undefined value. To prove this, we use the `Znth_map` lemma to move the `Znth` inside the `Vint`, leaving the goal,

```
is_int I32 Signed (Vint (Znth i (map Int.repr contents) Int.zero))
```

This is an instance of `is_int I32 Signed (Vint ...)` which is $\beta\eta$ -equal to `True`. However, when we rewrote by `Znth_map`, that left a subgoal,

```
HRE : i < Zlength (map Vint (map Int.repr contents))
H1 : 0 ≤ i ≤ Zlength (map Vint (map Int.repr contents))
  (other above-the-line hypotheses elided)
────────────────────────────────────────────────────────────────────────────────
0 ≤ i < Zlength (map Int.repr contents)
```

This solves straightforwardly as shown in the proof script.

12 Splitting sublists

In `progs/verif_sumarray.v`, at the comment “Now we have reached the end of the loop body,” it is time to prove that the *current* precondition (which is the postcondition of the loop body) entails the loop invariant. This is the proof goal:

```

H : 0 ≤ size ≤ Int.max_signed
H0 : Forall (fun x : Z ⇒ Int.min_signed ≤ x ≤ Int.max_signed) contents
HRE : i < size
H1 : 0 ≤ i ≤ size
  (other above-the-line hypotheses elided)


---


ENTAIL Delta,
PROP()
LOCAL(temp _i (Vint (Int.add (Int.repr i) (Int.repr 1))));
temp _s
  (force_val
    (sem_add_default tint tint
      (Vint (Int.repr (sum_Z (sublist 0 i contents)))))
      (Znth i (map Vint (map Int.repr contents)) Vundef)));
temp _x (Znth i (map Vint (map Int.repr contents)) Vundef); temp _a a;
temp _n (Vint (Int.repr size)))
SEP(data_at sh (tarray tint size) (map Vint (map Int.repr contents)) a)
⊢ EX a0 : Z,
  PROP(0 ≤ a0 ≤ size)
  LOCAL(temp _a a; temp _i (Vint (Int.repr a0)));
  temp _n (Vint (Int.repr size));
  temp _s (Vint (Int.repr (sum_Z (sublist 0 a0 contents)))))
  SEP(data_at sh (tarray tint size) (map Vint (map Int.repr contents)) a)

```

The right-hand side of this entailment is just the loop invariant. As usual at the end of a loop body, there is an existentially quantified variable that must be instantiated with an iteration-dependent value. In this case it’s obvious: the quantified variable represents the contents of C local variable `_i`, so we do, **Exists** (i+1).

The resulting entailment has many trivial parts and a nontrivial residue. The usual way to get to the hard part is to run `entailer!`, which we do now. After clearing away the irrelevant hypotheses, we have:

$$\begin{array}{l}
 H : 0 \leq \text{Zlength } (\text{map } \text{Vint } (\text{map } \text{Int.repr } \text{contents})) \leq \text{Int.max_signed} \\
 HRE : i < \text{Zlength } (\text{map } \text{Vint } (\text{map } \text{Int.repr } \text{contents})) \\
 H1 : 0 \leq i \leq \text{Zlength } (\text{map } \text{Vint } (\text{map } \text{Int.repr } \text{contents})) \\
 \text{-----}(1/1) \\
 \text{Vint } (\text{Int.repr } (\text{sum_Z } (\text{sublist } 0 \ (i + 1) \ \text{contents}))) = \\
 \text{force_val} \\
 (\text{sem_add_default tint tint } (\text{Vint } (\text{Int.repr } (\text{sum_Z } (\text{sublist } 0 \ i \ \text{contents})))) \\
 (\text{Znth } i \ (\text{map } \text{Vint } (\text{map } \text{Int.repr } \text{contents})) \ \text{Vundef}))
 \end{array}$$

The `sem_add_default` comes from the semantics of C expression evaluation: adding integers means one thing, but adding an integer to a `Vundef` is undefined, and so on. To clear that sludge out of the way, we move the `Znth` inside the `Vint` just as on [page 27](#), then `simpl`, yielding this goal:

$$\begin{array}{l}
 H : 0 \leq \text{Zlength } \text{contents} \leq \text{Int.max_signed} \\
 HRE : i < \text{Zlength } \text{contents} \\
 H1 : 0 \leq i \leq \text{Zlength } \text{contents} \\
 \text{-----}(1/1) \\
 \text{Vint } (\text{Int.repr } (\text{sum_Z } (\text{sublist } 0 \ (i + 1) \ \text{contents}))) = \\
 \text{Vint } (\text{Int.add } (\text{Int.repr } (\text{sum_Z } (\text{sublist } 0 \ i \ \text{contents}))) \\
 (\text{Int.repr } (\text{Znth } i \ \text{contents } 0)))
 \end{array}$$

The lemma `add_repr: $\forall i \ j, \text{Int.add } (\text{Int.repr } i) (\text{Int.repr } j) = \text{Int.repr } (i + j)$` is useful here; followed by `f_equal`, leaves:

$$\begin{array}{l}
 \text{sum_Z } (\text{sublist } 0 \ (i + 1) \ \text{contents}) = \\
 \text{sum_Z } (\text{sublist } 0 \ i \ \text{contents}) + \text{Znth } i \ \text{contents } 0
 \end{array}$$

Now the lemma `sublist_split: $\forall l \ m \ h \ al, \ 0 \leq l \leq m \leq h \leq |al| \rightarrow \text{sublist } l \ h \ al = \text{sublist } l \ m \ al ++ \text{sublist } m \ h \ al$` is helpful here: rewrite `(sublist_split 0 i (i+1))` by `omega`. A bit more rewriting with the theory of `sum_Z` and `sublist` finishes the proof.

13 Returning from a function

In `progs/verif_sumarray.v`, at the comment “After the loop,” we have reached the return statement. The forward tactic works here, leaving a proof goal that the precondition of the return entails the postcondition of the function-spec. (When this automatically, it leaves no proof goal at all.) The goal is a *lowered* entailment (on `mpred` assertions).

After doing `simpl` to clear away some C-expression-evaluation sludge, we have

```
H4 : Forall (value_fits tint) (map Vint (map Int.repr contents))
H2 : field_compatible (Tarray tint (Zlength ...) noattr) [] a
  (other above-the-line hypotheses elided)
-----
data_at sh (tarray tint (Zlength ...)) (map Vint (map Int.repr contents)) a
⊢ !!(Vint (Int.repr (sum_Z contents)) =
    Vint (Int.repr (sum_Z (sublist 0 i contents))))
```

The left-hand side of this entailment is a spatial predicate (`data_at`). Purely nonspatial facts (`H4` and `H2`) derivable from it have already been inferred and moved above the line by `saturate_local` (see [Chapter 30](#)).

This entailment’s right-hand side has no spatial predicates. That’s because the SEP clause of the funspec’s postcondition had exactly the same `data_at` clause as we see here in the entailment precondition, and the entailment-solver called by `forward` has already cleared it away.

In a situation like this—where `saturate_local` has already been done *and* the r.h.s. of the entailment is purely nonspatial—*almost always* there’s no more useful information in the left hand side that hasn’t already been extracted by `saturate_local`. We can throw away the l.h.s. with `apply prop_right` (or by `entailer!` but that’s a bit slower).

The remaining subgoal solves easily in the theory of sublists. The proof of the function `sumarray` is now complete.

14 *Global variables and* `main()`

C programs may have “extern” global variables, either with explicit initializers or initialized by default. Any function that accesses a global variable must have the appropriate spatial assertions in its funspec’s precondition (and postcondition). But the main function is special: it has spatial assertions for *all* the global variables. Then it may pass these on, piecemeal, to the functions it calls on an as-needed basis.

The function-spec for main always looks the same:

Definition `main_spec` :=

```
DECLARE _main WITH u : unit
  PRE [] main_pre prog u
  POST [ tint ] main_post prog u.
```

`main_pre` calculates the precondition automatically from (the list of extern global variables and initializers of) the program. Then, when we prove that main satisfies its funspec,

Lemma `body_main`: `semax_body Vprog Gprog f_main main_spec`.

Proof.

```
name four _four.
start_function.
```

the `start_function` tactic “unpacks” `main_pre` into an assertion:

```
four : val
----- (1/1)
semax Delta
  (PROP () LOCAL(gvar _four four)
    SEP(data_at Ews (tarray tint 4)
      (map Vint [Int.repr 1; Int.repr 2; Int.repr 3; Int.repr 4]) four))
  (... function body ...)
POSTCONDITION
```

The `LOCAL` clause means that the C global variable `_four` is at memory address *four*. (If we had omitted the name tactic in the proof script above, then `start.funcon` would have chosen some other name for this value.) See [Chapter 28](#).

The `SEP` clause means that there's data of type “array of 4 integers” at address *four*, with access permission `Ews` and contents `[1;2;3;4]`. `Ews` stands for “external write share,” the standard access permission of extern global writable variables. See [Chapter 35](#).

Now it's time to prove the function-call statement, `s = sumarray(four,4)`. When proving a function call, one must supply a *witness* for the `WITH` clause of the function-spec. The `_sumarray` function's `WITH` clause binds variables `a:val`, `sh:share`, `contents:list Z`, `size: Z`, so the type of the witness will be `(val*(share*(list Z * list Z)))`. To choose the witness, examine your actual parameter values (along with the precondition of the funspec) to see what witness would be consistent; here, we use `(four,Ews,four_contents,4)`.
`forward_call (four,Ews,four_contents,4).`

The `forward_call` tactic (usually) leaves subgoals: you must prove that your current precondition implies the funspec's precondition. Here, these solve easily, as shown in the proof script.

The postcondition of the call statement (which is the precondition of the next return statement) has an existential, `EX vret:val`. This comes directly from the existential in the funspec's postcondition. To move `vret` above the line, simply `Intros vret`.

Finally, we are at the return statement. The `forward` tactic is easily able to prove that the current assertion implies the postcondition of `_main`, because `main_post` is basically an abbreviation for `True`.

15 Tying all the functions together

We build a whole-program proof by composing together the proofs of all the function bodies. Consider `Gprog`, the list of all the function-specifications:

Definition `Gprog : funspecs := sumarray.spec :: main.spec::nil`.

Each `semax.body` proof says, assuming that *all the functions I might call* behave as specified, then *my own function-body* indeed behaves as specified:

Lemma `body_sumarray: semax.body Vprog Gprog f_sumarray sumarray.spec`.

Note that *all the functions I might call* might even include “myself,” in the case of a recursive or mutually recursive function.

This might seem like circular reasoning, but it is actually sound—by the miracle of step-indexed semantic models, as explained in Chapters 18 and 39 of *Program Logics for Certified Compilers*.

The rule for tying the functions together is called `semax.func`, and its use is illustrated in this theorem, the main proof-of-correctness theorem for the program `sumarray.c`:

Lemma `all_funcs_correct: semax.func Vprog Gprog (prog_func prog) Gprog`.
Proof.

`unfold Gprog, prog, prog_func; simpl.`

`semax.func_skipn.`

`semax.func_cons body_sumarray.`

`semax.func_cons body_main.`

`apply semax.func_nil.`

Qed.

The calls to `semax.func_cons` must appear in the same order as the functions are listed in `Gprog` and the same order as they appear in `prog.(prog_defs)`.

16 Separation logic: EX , $*$, emp , $!!$

The *base level* separation logic is built, like any separation logic, from predicates on “heaplets”. The grammar of base-level separation-logic expressions is,

$R ::=$	emp	empty
	$R_1 * R_2$	separating conjunction
	$R_1 \&\& R_2$	ordinary conjunction
	$\text{field_at } \pi \ \tau \ \vec{fld} \ v \ p$	“field maps-to”
	$\text{data_at } \pi \ \tau \ v \ p$	“maps-to”
	$\text{array_at } \tau \ \pi \ v \ lo \ hi$	array slice
	$!!P$	pure proposition
	$\text{EX } x : T, R$	existential quantification
	$\text{ALL } x : T, R$	universal quantification (rare)
	$R_1 \parallel R_2$	disjunction
	$\text{wand } R \ R'$	magic wand $R \multimap R'$ (rare)
	\dots	other operators, including user definitions

17 PROP() LOCAL() SEP()

The *lifted* separation logic can “see” local and global variables of the C program, in addition to the contents of the heap (pointer dereferences) that the base level separation logic can see. The *canonical form* of a lifted assertion is $\text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(\vec{R})$, where \vec{P} is a list of propositions (Prop), where \vec{Q} is a list of local-variable definitions (localdef), and \vec{R} is a list of base-level assertions (mpred). Each list is semicolon-separated.

Lifted assertions can occur in other forms than canonical form; in fact, anything of type $\text{environ} \rightarrow \text{mpred}$ is a lifted assertion. But canonical form is most convenient for forward symbolic execution (Hoare-logic rules).

The existential quantifier EX can also be used on canonical forms, e.g., $\text{EX } x:T, \text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(\vec{R})$.

Entailments in canonical form are normally of the form, $\text{ENTAIL } \Delta, PQR \vdash PQR'$, where PQR is a lifted assertion in canonical form, PQR' is a lifted assertion not necessarily in canonical form, and Δ is a type context. The \vdash operator is written $|-$ in Coq.

This notation is equivalent to $(\text{tc_environ } \Delta \ \&\& \ PQR) \vdash PQR'$. That is, Δ just provides extra assertions on the left-hand side of the entailment.

18 EX, Intros, Exists

In a canonical-form lifted assertion, existentials can occur at the outside, or in one of the base-level conjuncts within the SEP clause. This assertion has both:

```

ENTAIL  $\Delta$ ,
  EX  $x:Z$ ,
    PROP( $0 \leq x$ ) LOCAL(temp _i (Vint (Int.repr x)))
    SEP(EX  $y:Z$ ,  $!!(x < y) \ \&\& \text{data\_at } \pi \text{ tint (Vint (Int.repr y)) } p$ )
 $\vdash$  EX  $u: Z$ ,
  PROP( $0 < u$ ) LOCAL()
  SEP(data_at  $\pi$  tint (Vint (Int.repr  $u$ ))  $p$ )

```

To prove this entailment, one can first move x and y “above the line” by the tactic **Intros** a b:

```

 $a: Z$ 
 $b: Z$ 
 $H: 0 \leq a$ 
 $H0: a < b$ 

```

```

ENTAIL  $\Delta$ ,
  PROP() LOCAL(temp _i (Vint (Int.repr  $a$ )))
  SEP(data_at  $\pi$  tint (Vint (Int.repr  $b$ ))  $p$ )
 $\vdash$  EX  $u: Z$ ,
  PROP( $0 < u$ ) LOCAL()
  SEP(data_at  $\pi$  tint (Vint (Int.repr  $u$ ))  $p$ )

```

One might just as well say **Intros** $x \ y$ to use those names instead of a b. Note that the propositions (previously hidden inside existential quantifiers) have been moved above the line by **Intros**. Also, if there had been any separating-conjunction operators $*$ within the SEP clause, those will be “flattened” into semicolon-separated conjuncts within SEP.

Sometimes, even when there are no existentials to introduce, one wants to move PROP propositions above the line and flatten the $*$ operators into semicolons. One can just say **Intros** with no arguments to do that.

If you want to Intro an existential *without* gratuitous PROP-introduction and $*$ -flattening, you can just use **Intro** a , instead of **Intros** a .

Then, instantiate u by **Exists** b .

$a: Z$

$b: Z$

H: $0 \leq a$

H0: $a < b$

ENTAIL Δ ,
 PROP() LOCAL(temp _i (Vint (Int.repr a)))
 SEP(data_at π tint (Vint (Int.repr b)) p)
 \vdash PROP($0 < b$) LOCAL()
 SEP(data_at π tint (Vint (Int.repr b)) p)

This entailment proves straightforwardly by entailer!.

19 Integers: nat, Z, int 38 (compcert/lib/Integers.v)

Coq's standard library has the natural numbers `nat` and the integers `Z`. C-language integer values are represented by the type `Int.int` (or just `int` for short), which are 32-bit two's complement signed or unsigned integers with $\text{mod-}2^{32}$ arithmetic.

For most purposes, specifications and proofs of C programs should use `Z` instead of `nat`. Subtraction doesn't work well on naturals, and that screws up many other kinds of arithmetic reasoning. *Only when you are doing direct natural-number induction* is it natural to use `nat`, and so you might then convert using `Z.to_nat` to do that induction.

Conversions between `Z` and `int` are done as follows:

`Int.repr`: $Z \rightarrow \text{int}$.

`Int.unsigned`: $\text{int} \rightarrow Z$.

`Int.signed`: $\text{int} \rightarrow Z$.

with the following lemmas:

$$\text{Int.repr_unsigned} \frac{}{\text{Int.repr}(\text{Int.unsigned } z) = z}$$

$$\text{Int.unsigned_repr} \frac{0 \leq z \leq \text{Int.max_unsigned}}{\text{Int.unsigned}(\text{Int.repr } z) = z}$$

$$\text{Int.repr_signed} \frac{}{\text{Int.repr}(\text{Int.signed } z) = z}$$

$$\text{Int.signed_repr} \frac{\text{Int.min_signed} \leq z \leq \text{Int.max_signed}}{\text{Int.signed}(\text{Int.repr } z) = z}$$

`Int.repr` truncates to a 32-bit twos-complement representation (losing information if the input is out of range). `Int.signed` and `Int.unsigned` are different injections back to `Z` that never lose information.

When doing proofs about integers, the recommended proof technique is to make sure your integers never overflow. That is, if the C variable `_x`

contains the value `Vint (Int.repr x)`, then make sure x is in the appropriate range. Let's assume that x is a signed integer, i.e. declared in C as `int x`; then the hypothesis is,

$H: \text{Int.min_signed} \leq x \leq \text{Int.max_signed}$

If you maintain this hypothesis “above the line”, then Floyd’s tactical proof automation can solve goals such as `Int.signed (Int.repr x) = x`. Also, to solve goals such as,

...

$H2 : 0 \leq n \leq \text{Int.max_signed}$

...

$\text{Int.min_signed} \leq 0 \leq n$

you can use the `reapable_signed` tactic, which is basically just `omega` with knowledge of the values of `Int.min_signed`, `Int.max_signed`, and `Int.max_unsigned`.

To take advantage of this, put conjuncts into the `PROP` part of your function precondition such as $0 \leq i < n; n \leq \text{Int.max_signed}$. Then the `start_function` tactic will move them above the line, and the other tactics mentioned above will make use of them.

To see an example in action, look at `progs/verif_sumarray.v`. The array size and index (variables `size` and `i`) are kept within bounds; but the *contents* of the array might overflow when added up, which is why `add_elem` uses `Int.add` instead of `Z.add`.

20 Values: Vint, Vptr (compcert/common/Values.v)

Definition block : Type := positive.

Inductive val: Type :=

| Vundef: val
 | Vint: int → val
 | Vlong: int64 → val
 | Vfloat: float → val
 | Vsingle: float32 → val
 | Vptr: block → int → val.

Vundef is the *undefined* value—found, for example, in an uninitialized local variable.

Vint(i) is an integer value, where i is a CompCert 32-bit integer. These 32-bit integers can also represent short (16-bit) and char (8-bit) values.

Vfloat(f) is a 64-bit floating-point value.

Vsingle(f) is a 32-bit floating-point value.

Vptr $b\ z$ is a pointer value, where b is an abstract block number and z is an offset within that block. Different *malloc* operations, or different extern global variables, or stack-memory-resident local variables, will have different abstract block numbers. Pointer arithmetic must be done within the same abstract block, with $(\text{Vptr } b\ z) + (\text{Vint } i) = \text{Vptr } b\ (z + i)$. Of course, the C-language $+$ operator first multiplies i by the size of the array-element that $\text{Vptr } b\ z$ points to.

Vundef is not always treated as distinct from a defined value. For example, $p \mapsto \text{Vint } 5 \vdash p \mapsto \text{Vundef}$, where \mapsto is the `data_at` operator ([Chapter 25](#)). That is, $p \mapsto \text{Vundef}$ really means $\exists v, p \mapsto v$. Vundef could mean “truly uninitialized” or it could mean “initialized but arbitrary.”

CompCert C describes C's type system with inductive data types.

Inductive signedness := Signed | Unsigned.

Inductive intsize := l8 | l16 | l32 | lBool.

Inductive floatsize := F32 | F64.

Record attr : Type := mk_attr {
 attr_volatile: bool; attr_alignas: option N
}.

Definition noattr := {| attr_volatile := false; attr_alignas := None |}.

Inductive type : Type :=

| Tvoid: type
| Tint: intsize → signedness → attr → type
| Tlong: signedness → attr → type
| Tfloat: floatsize → attr → type
| Tpointer: type → attr → type
| Tarray: type → Z → attr → type
| Tfunction: typelist → type → calling_convention → type
| Tstruct: ident → attr → type
| Tunion: ident → attr → type

with typelist : Type :=

| Tnil: typelist
| Tcons: type → typelist → typelist.

We have abbreviations for commonly used types:

Definition tint = Tint l32 Signed noattr.

Definition tuint = Tint l32 Unsigned noattr.

Definition tschar = Tint l8 Signed noattr.

Definition tuchar = Tint l8 Unsigned noattr.

Definition tarray (t: type) (n: Z) = Tarray t n noattr.

Definition tptr (t: type) := Tpointer t noattr.

22 *CompSpecs*

The C language has a namespace for struct- and union-identifiers, that is, *composite types*. In this example, struct foo {int value; struct foo *tail} a,b; the “global variables” namespace contains a,b, and the “struct and union” namespace contains foo.

When you use CompCert clightgen to parse myprogram.c into myprogram.v, the main definition it produces is prog, the AST of the entire C program:

Definition prog : Clight.program := {| prog_types := composites; ... |}.

To interpret the meaning of a type expression, we need to look up the names of its struct identifiers in a *composite* environment. This environment, along with various well-formedness theorems about it, is built from prog as follows:

Require Import floyd.proofauto. (** Import Verifiable C library **)

Require Import myprogram. (** AST of my program **)

Instance CompSpecs : compspecs. **Proof.** make_compspecs prog. **Defined.**

The make_compspecs tactic automatically constructs the *composite specifications* from the program. As a typeclass Instance, CompSpecs is supplied automatically as an implicit argument to the functions and predicates that interpret the meaning of types:

Definition sizeof {env: composite.env} (t: type) : Z := ...

Definition data_at_ {cs: compspecs} (sh: share) (t: type) (v: val) := ...

@sizeof (@cenv.cs CompSpecs) (Tint l32 Signed noattr) = 4.

sizeof (Tint l32 Signed noattr) = 4.

sizeof (Tstruct _foo noattr) = 8.

@data_at_ CompSpecs sh t v ⊢ data_at_ sh t v

When you have two separately compiled .c files, each will have its own prog and its own compspecs. See [Chapter 38](#).

23 retype

For each C-language data type, we define a *representation type*, the Type of Coq values that represent the contents of a C variable of that type.

Definition `retype {cs: compspecs} (t: type) : Type := ...`

Lemma `retype_ind: $\forall (t: \text{type}),$`

`retype t =`

`match t with`

`| Tvoid \Rightarrow unit`

`| Tint _ _ \Rightarrow val`

`| Tlong _ _ \Rightarrow val`

`| Tfloat _ _ \Rightarrow val`

`| Tpointer _ _ \Rightarrow val`

`| Tarray t0 _ _ \Rightarrow list (retype t0)`

`| Tfunction _ _ _ \Rightarrow unit`

`| Tstruct id _ \Rightarrow retype_structlist (co_members (get_co id))`

`| Tunion id _ \Rightarrow retype_unionlist (co_members (get_co id))`

`end`

`retype_structlist` is the right-associative cartesian product of all the (retypes of) the fields of the struct. For example,

`struct list {int hd; struct list *tl};`

`struct one {struct list *p};`

`struct three {int a; struct list *p; double x};`

`retype (Tstruct _list noattr) = (val*val).`

`retype (Tstruct _one noattr) = val.`

`retype (Tstruct _three noattr) = (val*(val*val)).`

We use `val` instead of `int` for the retype of an integer variable, because the variable might be uninitialized, in which case its value will be `Vundef`.

24 *Uninitialized data*, default_val

CompCert represents uninitialized atomic (integer, pointer, float) values as `Vundef` : `val`.

The dependently typed function `default_val` calculates the undefined value for any C type:

`default_val`: $\forall \{cs: \text{compspecs}\} (t: \text{type}), \text{reptype } t.$

For any C type t , the default value for variables of type t will have Coq type `(reptype t)`.

For example:

```
struct list {int hd; struct list *tl;};
```

```
default_val tint = Vundef
```

```
default_val (tptr tint) = Vundef
```

```
default_val (tarray tint 4) = [Vundef; Vundef; Vundef; Vundef]
```

```
default_val (tarray  $t$   $n$ ) = list_repeat (Z.to_nat  $n$ ) (default_val  $t$ )
```

```
default_val (Tstruct _list noattr) = (Vundef, Vundef)
```

25 data_at

Consider a C program with these declarations:

```
struct list {int hd; struct list *tl;} L;
int f(struct list a[5], struct list *p) { ... }
```

Assume these definitions in Coq:

Definition t_list := Tstruct _list noattr.

Definition t_arr := Tarray t_list 5 noattr.

Somewhere inside `f`, we might have the assertion,

```
PROP() LOCAL(temp _a  $\alpha$ , temp _p  $p$ , gvar _L  $L$ )
SEP(data_at Ews t_list (Vint (Int.repr 0), nullval)  $L$ ;
    data_at  $\pi$  t_arr (list_repeat (Z.to_nat 5) (Vint (Int.repr 1),  $p$ ))  $\alpha$ ;
    data_at  $\pi$  t_list (default_val t_list)  $p$ )
```

This assertion says, “Local variable `_a` contains address α , `_p` contains address p , global variable `_L` is at address L . There is a struct list at L with permission-share `Ews` (“extern writable share”), whose `hd` field contains 0 and whose `tl` contains a null pointer. At address α there is an array of 5 list structs, each with `hd`=1 and `tl`= p , with permission π ; and at address p there is a single list cell that is uninitialized¹, with permission π .”

In pencil-and-paper separation logic, we write $q \mapsto i$ to mean `data_at Tsh tint (Vint (Int.repr i)) q` . We write $L \mapsto (0, \text{NULL})$ to mean `data_at Tsh t_list (Vint (Int.repr 0), nullval) L` . We write $p \mapsto (_, _)$ to mean `data_at π t_list (default_val t_list) p` .

In fact, the definition `data_at_` is useful for the situation $p \mapsto _$:

Definition data_at_ {cs: compspecs} sh t p := data_at sh t (default_val t) p.

¹Uninitialized, or initialized but we don’t know or don’t care what its value is

26 retype', repinj

```

struct a {double x1; int x2;};
struct b {int y1; struct a y2;} p;
repinj:  $\forall t$ : type, retype' t  $\rightarrow$  retype t
retype t_struct_b = (val*(val*val))
retype' t_struct_b = (int*(float*int))
repinj t_struct_b (i,(x,j)) = (Vint i, (Vfloat x, Vint j))

```

TL;DR

The retype function maps C types to the the corresponding Coq types of (possibly uninitialized) values. When we know a variable is definitely initialized, it may be more natural to use int instead of val for integer variables, and float instead of val for double variables. The retype' function maps C types to the Coq types of (definitely initialized) values.

Definition retype' {cs: compspecs} (t: type) : Type := ...

Lemma retype'_ind: $\forall (t$: type),
retype t =

```

match t with
| Tvoid  $\Rightarrow$  unit
| Tint _ _  $\Rightarrow$  int
| Tlong _ _  $\Rightarrow$  Int64.int
| Tfloat _ _  $\Rightarrow$  float
| Tpointer _ _  $\Rightarrow$  pointer_val
| Tarray t0 _ _  $\Rightarrow$  list (retype' t0)
| Tfunction _ _ _  $\Rightarrow$  unit
| Tstruct id _  $\Rightarrow$  retype'_structlist (co_members (get_co id))
| Tunion id _  $\Rightarrow$  retype'_unionlist (co_members (get_co id))
end

```

The function repinj maps an initialized value to the type of possibly uninitialized values:

Definition repinj {cs: compspecs} (t: type) : retype' t \rightarrow retype t := ...

The program `progs/nest2.c` (verified in `progs/verif_nest2.v`) illustrates the use of `retype'` and `repinj`.

```
struct a {double x1; int x2;};
struct b {int y1; struct a y2;} p;

int get(void) { int i; i = p.y2.x2; return i; }
void set(int i) { p.y2.x2 = i; }
```

Our API spec for `get` reads as,

Definition `get_spec` :=

```
DECLARE _get
  WITH v : retype' t_struct_b, p : val
  PRE []
    PROP() LOCAL(gvar _p p)
    SEP(data_at Ews t_struct_b (repinj _ v) p)
  POST [ tint ]
    PROP() LOCAL(temp ret_temp (Vint (snd (snd v))))
    SEP(data_at Ews t_struct_b (repinj _ v) p).
```

In this program, `retype' t_struct_b = (int*(float*int))`, and `repinj t_struct_b (i,(x,j)) = (Vint i, (Vfloat x, Vint j))`.

One could also have specified `get` without `retype'` at all:

Definition `get_spec` :=

```
DECLARE _get
  WITH i: Z, x: float, j: int, p : val
  PRE []
    PROP() LOCAL(gvar _p p)
    SEP(data_at Ews t_struct_b (Vint (Int.repr i), (Vfloat x, Vint j)) p)
  POST [ tint ]
    PROP() LOCAL(temp ret_temp (Vint j))
    SEP(data_at Ews t_struct_b (Vint (Int.repr i), (Vfloat x, Vint j)) p).
```

27 field_at

Consider again the example in `progs/nest2.c`

```
struct a {double x1; int x2;};
struct b {int y1; struct a y2;};
```

The command `i = p.y2.x2;` does a nested field load. We call `y2.x2` the *field path*. The precondition for this command might include the assertion,

```
LOCAL(gvar _pb pb)
SEP( data_at sh t_struct_b (y1,(x1,x2)) pb)
```

The postcondition (after the load) would include the new `LOCALfact`,
`temp _i x2.`

The tactic (`unfold_data_at 1%nat`) changes the `SEP` part of the assertion as follows:

```
SEP(field_at Ews t_struct_b (DOT _y1) (Vint y1) pb;
   field_at Ews t_struct_b (DOT _y2) (Vfloat x1, Vint x2) pb)
```

and then doing (`unfold_field_at 2%nat`) unfolds the second `field_at`,

```
SEP(field_at Ews t_struct_b (DOT _y1) (Vint y1) pb;
   field_at Ews t_struct_b (DOT _y2 DOT _x1) (Vfloat x1) pb;
   field_at Ews t_struct_b (DOT _y2 DOT _x2) (Vint x2) pb)
```

The third argument of `field_at` represents the *path* of structure-fields that leads to a given substructure. The empty path (`nil`) works too; it “leads” to the entire structure. In fact, `data_at $\pi \tau v p$` is just short for `field_at $\pi \tau nil v p$` .

Arrays and structs may be nested together, in which case the field path may also contain array subscripts at the appropriate places, using the notation `SUB i` along with `DOT field`.

28 *Localdefs*: temp, lvar, gvar

The `LOCAL` part of a `PROP()``LOCAL()``SEP()` assertion is a list of *localdefs* that bind variables to their values or addresses.

Inductive `localdef` : `Type` :=
 | `temp`: `ident` → `val` → `localdef`
 | `lvar`: `ident` → `type` → `val` → `localdef`
 | `gvar`: `ident` → `val` → `localdef`
 | `sgvar`: `ident` → `val` → `localdef`
 | `localprop`: `Prop` → `localdef`.

`temp i v` binds a nonaddressable local variable *i* to its value *v*.

`lvar i t v` binds an *addressable* local variable *i* (of type *t*) to its *address* *v*.

`gvar i v` binds a *visible global* variable *i* to its *address* *v*.

`sgvar i v` binds a *possibly shadowed global* variable *i* to its *address* *v*.

The *contents* of an addressable (local or global) variable is on the heap, and can be described in the `SEP` clause.

```
int g=2;
int f(void) { int g; int *p = &g; g=6; return g; }
```

In this program, the global variable *g* is shadowed by the local variable *g*. In an assertion inside the function body, one could write

```
PROP() LOCAL(temp _p q; lvar _g tint q; sgvar _g p}
SEP(data_at Ews tint (Vint (Int.repr 2)) p; data_at Tsh tint (Vint (Int.repr 6)) q)
```

to describe a shadowed global variable *_g* that is still there in memory but (temporarily) cannot be referred to by its name in the C program.

29 *go_lower*

Normally one does not use this tactic directly, it is invoked as the first step of *entailer* or *entailer!*

Given a lifted entailment $\text{ENTAIL } \Delta, \text{PROP}(\vec{P}) \text{ LOCAL}(\vec{Q}) \text{ SEP}(\vec{R}) \vdash S$, one often wants to prove it at the base level: that is, with all of \vec{P} moved above the line, with all of \vec{Q} out of the way, just considering the base-level separation-logic conjuncts \vec{R} .

When $\Delta, \vec{P}, \vec{Q}, \vec{R}$ are *concrete*, the *go_lower* tactic does this. Concrete means that the \vec{P}, \vec{Q} are nil-terminated lists (not Coq variables) that every element of \vec{Q} is manifestly a *localdef* (not hidden in Coq abstractions), the identifiers in \vec{Q} be (computable to) ground terms, and the analogous (tree) property for Δ . It is not necessary that $\Delta, \vec{P}, \vec{Q}, \vec{R}$ be fully *ground terms*: Coq variables (and other Coq abstractions) can appear anywhere in \vec{P} and \vec{R} and in the *value* parts of Δ and \vec{Q} . When the entailment is not fully concrete, or when there existential quantifiers outside *PROP*, the tactic *old_go_lower* can still be useful.

go_lower moves the propositions \vec{P} above the line; when a proposition is an equality on a Coq variable, substitute the variable.

For each *localdef* in \vec{Q} (such as *temp i v*), *go_lower* looks up *i* in Δ to derive a type-checking fact (such as *tc_val t v*), then introduces it above the line and simplifies it. For example, if *t* is *tptr tint*, then the typechecking fact simplifies to *is_pointer_or_null v*.

Then it proves the *localdefs* in *S*, if possible. If there are still some local-environment dependencies remaining in *S*, it introduces a variable *rho* to stand for the run-time environment.

The remaining goal will be of the form $\vec{R} \vdash S'$, with the semicolons in \vec{R} replaced by the separating conjunction ***. *S'* is the residue of *S* after lowering to the base separation logic and deleting its (provable) *localdefs*.

30 *saturate_local*

Normally one does not use this tactic directly, it is invoked by *entailer* or *entailer!*

To prove an entailment $R_1 * R_2 * \dots * R_n \vdash!! (P'_1 \wedge \dots \wedge P'_n) \&\& R'_1 * \dots * R'_m$, first extract all the *local (nonspatial)* facts from $R_1 * R_2 * \dots * R_n$, use them (along with other propositions above the line) to prove $P'_1 \wedge \dots \wedge P'_n$, and then work on the separation-logic (spatial) conjuncts $R_1 * \dots * R_n \vdash R'_1 * \dots * R'_m$.

An example local fact: $\text{data_at } \text{Ews } (\text{tarray tint } n) \ v \ p \vdash!! (\text{Zlength } v = n)$. That is, the value v in an array “fits” the length of the array.

The Hint database *saturate_local* contains all the local facts that can be extracted from *individual* spatial conjuncts:

field_at_local_facts:

$$\begin{aligned} \text{field_at } \pi \ t \ \text{path} \ v \ p \vdash!! & (\text{field_compatible } t \ \text{path} \ p \\ & \wedge \text{value_fits } (\text{nested_field_type } t \ \text{path}) \ v) \\ \text{data_at } \pi \ t \ v \ p \vdash!! & (\text{field_compatible } t \ \text{nil} \ p \wedge \text{value_fits } t \ v) \end{aligned}$$

memory_block_local_facts:

$$\text{memory_block } \pi \ n \ p \vdash!! \text{isptr } p$$

The assertion $(\text{Zlength } v = n)$ is actually a consequence of *value_fits* when t is an array type. See [Chapter 32](#).

If you create user-defined spatial terms (perhaps using EX, *data_at*, etc.), you can add hints to the *saturate_local* database as well.

The tactic *saturate_local* takes a proof goal of the form $R_1 * R_2 * \dots * R_n \vdash S$ and adds *saturate-local* facts for *each* of the R_i , though it avoids adding duplicate hypotheses above the line.

31 *field_compatible, field_address*

CompCert C light comes with an “address calculus.” Consider this example:

```
struct a {double x1; int x2;};
struct b {int y1; struct a y2;};
struct a *pa; int *q = &(pa→y2.x2);
```

Suppose the value of `_pa` is p . Then the value of `_q` is $p + \delta$; how can we reason about δ ?

Given type t such as `Tstruct _b noattr`, and $path$ such as `(DOT _y2 DOT _x2)`, then `(nested_field_type t path)` is the type of the field accessed by that path, in this case `tint`; `(nested_field_offset t path)` is the distance (in bytes) from the base of t to the address of the field, in this case (on a 32-bit machine) 12 or 16, depending on the field-alignment conventions of the target-machine.

On the Intel x86 architecture, where doubles need not be 8-byte-aligned, we have,

$$\text{data_at } \pi \text{ t_struct_b } (i, (f, j)) \ p \vdash \\ \text{data_at } \pi \text{ tint } i \ p * \text{data_at } \pi \text{ t_struct_a } (f, j) \ (\text{offset_val } p \ 12)$$

but don't write it that way! For one thing, the converse is not valid:

$$\text{data_at } \pi \text{ tint } i \ p * \text{data_at } \pi \text{ t_struct_a } (f, j) \ (\text{offset_val } p \ 12) \\ \not\vdash \text{data_at } \pi \text{ t_struct_b } (i, (f, j)) \ p$$

The reasons: we don't know that $p + 12$ satisfies the alignment requirements for struct b; we don't know whether $p + 12$ crosses the end-of-memory boundary. That entailment *would* be valid in the presence of this hypothesis: `field_compatible t_struct_b nil p : Prop`.

which says that an entire struct b value *can* fit at address p . Note that

this is a *nonspatial* assertion, a pure proposition, independent of the *contents* of memory.

In order to assist with reasoning about reassembly of data structures, `saturate_local` (and therefore `entailer`) put `field_compatible` assertions above the line; see [Chapter 30](#).

Sometimes one needs to name the address of an internal field—for example, to pass just that field to a function. In that case, one *could* use `field_offset`, but it better to use `field_address`:

Definition `field_address` (t : type) ($path$: list gfield) (p : val) : val :=
 if `field_compatible_dec t path p`
 then `offset_val (Int.repr (nested_field_offset t path)) p`
 else `Vundef`

That is, `field_address` has “baked in” the fact that the offset is “compatible” with the base address (is properly aligned, has not crossed the end-of-memory boundary). And therefore:

```
data_at  $\pi$  tint  $i$   $p$ 
  * data_at  $\pi$  t_struct_a ( $f, j$ ) (field_address t_struct_b (DOT _y2 DOT _x2)  $p$ )
 $\vdash$  data_at  $\pi$  t_struct_b ( $i, (f, j)$ )  $p$ 
```

32 *value_fits*

The spatial maps-to assertion, $\text{data_at } \pi \ t \ v \ p$, says that there's a value v in memory at address p , filling the data structure whose C type is t (with permission π). A corollary is $\text{value_fits } t \ v$: v is a value that actually *can* reside in such a C data structure.

Value_fits is a recursive, dependently typed relation that is easier described by its induction relation; here, we present a simplified version that assumes that all types t are not volatile:

```

value_fits  $t \ v = \text{tc\_val}' \ t \ v$    (when  $t$  is an integer, float, or pointer type)
value_fits (tarray  $t' \ n$ )  $v = (\text{Zlength } v = \text{Z.max } 0 \ n) \wedge \text{Forall } (\text{value\_fits } t') \ v$ 
value_fits (Tstruct  $i \ \text{noattr}$ )  $(v_1, (v_2, (\dots, v_n))) =$ 
    value_fits (field_type  $f_1 \ v_1$ )  $\wedge \dots \wedge \text{value\_fits (field\_type } f_n \ v_n)$ 
    (when the fields of struct  $i$  are  $f_1, \dots, f_n$ )

```

The predicate $\text{tc_val}'$ says,

Definition $\text{tc_val}' (t: \text{type}) (v: \text{val}) := v \neq \text{Vundef} \rightarrow \text{tc_val } t \ v$.

Definition $\text{tc_val } (t: \text{type}) (v: \text{val}) :=$

```

    match  $t$  with
    | Tvoid  $\Rightarrow$  False
    | Tint  $\text{sz sg } _ \Rightarrow \text{is\_int } \text{sz sg}$ 
    | Tlong  $_ _ \Rightarrow \text{is\_long}$ 
    | Tfloat F32  $_ \Rightarrow \text{is\_single}$ 
    | Tfloat F64  $_ \Rightarrow \text{is\_float}$ 
    | Tpointer  $_ _$  | Tarray  $_ _ _$  | Tfunction  $_ _ _ \Rightarrow \text{is\_pointer\_or\_null}$ 
    | Tstruct  $_ _$  | Tunion  $_ _ \Rightarrow \text{isptr}$ 
end

```

So, an atomic value (int, float, pointer) fits *either* when it is Vundef or when it type-checks. We permit Vundef to “fit,” in order to accommodate partially initialized data structures in C.

Given an entailment $(A_1 * A_2) * ((A_3 * A_4) * A_5) \vdash A'_4 * (A'_5 * A'_1) * (A'_3 * A'_2)$ for any associative-commutative rearrangement of the A_i , and where (for each i), A_i is $\beta\eta$ equivalent to A'_i , then the cancel tactic will solve the goal. When we say A_i is $\beta\eta$ equivalence to A'_i , that is equivalent to saying that (change (A_i) with (A'_i)) would succeed.

If the goal has the form $(A_1 * A_2) * ((A_3 * A_4) * A_5) \vdash (A'_4 * B_1 * A'_1) * B_2$ where there is only a partial match, then cancel will remove the matching conjuncts and leave a subgoal such as $A_2 * A_3 * A_5 \vdash B_1 * B_2$.

If the goal is $(A_1 * A_2) * ((A_3 * A_4) * A_5) \vdash A'_4 * \top * A'_1$, where some terms cancel and the rest can be absorbed into \top , then cancel will solve the goal.

If the goal has the form

$$\frac{F := ?224 : \text{list}(\text{environ} \rightarrow \text{mpred})}{(A_1 * A_2) * ((A_3 * A_4) * A_5) \vdash A'_4 * (\text{fold_right sepcon emp } F) * A'_1}$$

where F is a *frame* that is an abbreviation for an uninstantiated logical variable of type $\text{list}(\text{environ} \rightarrow \text{mpred})$, then the cancel tactic will perform *frame inference*: it will unfold the definition F , instantiate the variable (in this case, to $A_2 :: A_3 :: A_5 :: \text{nil}$), and solve the goal. The frame may have been created by $\text{evar}(F : \text{list}(\text{environ} \rightarrow \text{mpred}))$, as part of forward symbolic execution through a function call.


WARNING: cancel can turn a provable entailment into an unprovable entailment. Consider this:

$$\frac{A * C \vdash B * C}{A * D * C \vdash C * B * D}$$

This goal is provable by first rearranging to $(A * C) * D \vdash (B * C) * D$. But cancel may aggressively cancel C and D , leaving $A \vdash B$, which is not provable. You might wonder, what kind of crazy hypothesis is $A * C \vdash B * C$; but indeed such “context-dependent” cancellations do occur in the theory of linked lists; see ?? and PLCC Chapter 19.

34 entailer!

The mapsto operator (and related operators) take a *permission share*, expressing whether the mapsto grants read permission, write permission, or some other fractional permission.



The *top* share, written Tsh or Share.top, gives total permission: to deallocate any cells within the footprint of this mapsto, to read, to write.

Share.split Tsh = (Lsh, Rsh)	
Share.split Lsh = (a, a')	Share.split Rsh = (b, b')
$a' \oplus b = c$	$\text{lub}(c, \text{Rsh}) = a' \oplus \text{Rsh} = d$
$\forall sh. \text{writable_share } sh \rightarrow \text{readable_share } sh$	
writable_share Ews	readable_share b
writable_share d	readable_share c
writable_share Tsh	$\neg \text{readable_share Lsh}$

Any share may be split into a *left half* and a *right half*. The left and right of the top share are given distinguished names Lsh, Rsh.

The right-half share of the top share (or any share containing it such as *d*) is sufficient to grant *write permission* to the data: “the right share is the write share.” A thread of execution holding only Lsh—or subshares of it such as *a, a'*—can neither read or write the object, but such shares are not completely useless: holding any nonempty share prevents other threads from deallocating the object.

Any subshare of Rsh, in fact any share that overlaps Rsh, grants *read* permission to the object. Overlap can be tested using the glb (greatest

lower bound) operator.

Whenever $(\text{mapsto } sh \ t \ v \ w)$ holds, then the share sh must include at least a read share, thus this give permission to load memory at address v to get a value w of type t .

To make sure sh has enough permission to write (i.e., $Rsh \subset sh$, we can say $\text{writable_share } sh : \text{Prop}$.

Memory obtained from `malloc` comes with the top share Tsh . Writable extern global variables and stack-allocated addressable locals (which of course must not be deallocated) come with the “extern writable share” Ews which is equal to Rsh . Read-only globals come with a half-share of Rsh .

Sequential programs usually have little need of any shares except the Tsh and Ews . However, many function specifications can be parameterized over any share (example: [page ??](#)), and this sort of generalized specification makes the functions usable in more contexts.

In C it is undefined to test deallocated pointers for equality or inequalities, so the Hoare-logic rule for pointer comparison also requires some permission-share; see [page 112](#).

36 *Pointer comparisons*

37 Proving larg(ish) programs

When your program is not all in one .c file, see also [Chapter 38](#). Whether or not your program is all in one .c file, you can prove the individual function bodies in separate .v files. This uses less memory, and (on a multicore computer with parallel make) saves time. To do this, put your API spec (up to the construction of Gprog in one file; then each `semax_body` proof in a separate file that imports the API spec.

Extraction of subordinate semax-goals

To ease memory pressure and recompilation time, it is often advisable to partition the proof of a function into several lemmas. Any proof state whose goal is a `semax-term` can be extracted as a stand-alone statement by invoking tactic `semax_subcommand V G F`. The three arguments are as in the statement of surrounding `semax-body` lemma, i.e. are of type `varspecs`, `funspecs`, and `function`.

The subordinate tactic `mkConciseDelta V G F Delta` can also be invoked individually, to simply display the type context `Delta` in concise form, as the application of a sequence of initializations to the host function's `func_tycontext`.

The freezer

A distinguishing feature of separation logic is the frame rule, i.e. the ability to modularly verify a statement w.r.t. its minimal resource footprint. Unfortunately, being phrased in terms of the syntactic program structure, the standard frame rule does not easily interact with forward symbolic execution as implemented by the Floyd tactics (and many other systems), as these continuously rearrange the associativity of statement sequencing to peel off the redex of the next *forward*, and (purposely) hide the program continuation as the abbreviation `MORE_COMMANDS`.

Resolving this conflict, Floyd's *freezer* abstraction provides a means for

flexible framing, by implementing a veil that opaquely hides selected items of a SEP clause from non-symbolic treatment by non-freezer tactics.

The freezer abstraction consists of two main tactics, *freeze* $N\ F$ and *thaw* F , where $N : \text{list nat}$ and F is a user-supplied (fresh) Coq name. The result of applying *freeze* $[i_1; \dots; i_n]\ F$ to a semax goal is to remove items i_1, \dots, i_n from the precondition's SEP clause, inserting the item *FRZL* F at the head of the SEP list, and adding a hypothesis $F := \text{abbreviate}$ to Coq's proof context.

The term *FRZL* F participates symbolically in all non-freezer tactics just like any other SEP item, so can in particular be canceled, and included in a function call's frame. Unfolding a freezer is not tied to the associativity structure of program statements but can be achieved by invoking *thaw* F , which simply replaces *FRZL* F by the the list of F 's constituents. As multiple freezers can coexists and freezers can be arbitrarily nested, SEP-clauses R effectively contain forests of freezers, each constituent being thawable independently and freezer-level by freezer-level.

Wrapping single *forward* or *forward_call* commands in a freezer often speeds up the processing time noticably, as invocations of subordinate tactics *entailer*, *cancel*, etc. are supplied with smaller and more symbolic proof goals. In our experience, applying the freezer throughout the proof of an entire function body typically yields a speedup of about 30% on average with improvements of up to 55% in some cases, while also easing the memory pressure and freeing up valuable real estate on the user's screen.

A more invasive implementation of a freezer-like abstraction would refine the $\text{PROP}(P)\ \text{LOCAL}(Q)\ \text{SEP}(R)$ structure to terms of the form $\text{PROP}(P)\ \text{LOCAL}(Q)\ \text{SEP}(R)\ \text{FR}(H)$ where $H : \text{list mpred}$. Again, terms in H would be treated opaquely by all tactics, and freezing/thawing would correspond to transfer rules between R and H . In either case, forward symbolic execution is reconciled with the frame rule, and the use of the mechanism is sound engineering practice as documentation of

programmer's insight is combined with performance improvements.

38 *Separate compilation*, `semax_ext`

What to do when your program is spread over multiple `.c` files.

Code preparation

In order to separate the namespaces of multiple files compiled by Comp-Cert's `clightgen` tool, it is necessary to apply

```
python fix_clightgen.py file1.v ...fileN.v
```

The script reads in the named files, concisely renames variables etc by making up new positives, and writes the modified files back to the given names.

39 Appendix: catalog of major tactics / commands

Below is an alphabetic catalog of the major floyd tactics. In addition to short descriptions, the entries indicate whether a tactic (or tactic notation) is typically user-applied [u], primarily of internal use [i] or is expected to be used at development-time but unlikely to appear in a finished proof script [d]. We also mention major interdependencies between tactics, and their points of definition.

[ui] andp_left2 Used frequently in combination with *derives_refl* to clean up proofs states at the end of a loop or function, or after a *forward_seq*.

Subordinate tactics: ...

Superordinate tactics: ...

Defined in XXX.v

[u] cancel Main vehicle for resolving entailments between mpred lists. Performs symbolic cancellation, i.e. performs little inspection of the internals of mpreds. May turn a provable goal into an unprovable goal, by cancelling too aggressively. Occasionally needs to be finished off with a *derives_refl*, in case of syntactically not mtaching compspecs¹.

Subordinate tactics: ...

Superordinate tactics: ...

Defined in XXX.v

[u] derives_refl' (Actually, a lemma, not a tactic). Often followed by *f_equal* to solve, say, an entailment between two almost identical *data_at* predicates.

¹Do these cases still occur? Any other cases?

[u] entailer Vehicle for proving entailments at the mpred level or the assertion level. . .

Say sth about (non-)inclusion of saturate_locals, normalize, cancel?

Subordinate tactics: . . .

Superordinate tactics: . . .

Defined in XXX.v

[ui] derives_refl and derives_refl' Subordinate tactics: . . .

Superordinate tactics: . . .

Defined in XXX.v

[u] drop_LOCAL n where $n : nat$. Removes the n th entry of a the LOCAL block of a semax_statement's precondition².

Subordinate tactics: . . .

Superordinate tactics: . . .

Defined in XXX.v

[u] entailer! Vehicle for proving entailments at the mpred level or the assertion level. . . In contrast to *entailer*, *entailer!* may turn a provable entailment into an unprovable one, but is usually more efficient.

Say sth about (non-)inclusion of saturate_locals, normalize, cancel?

Subordinate tactics: . . .

Superordinate tactics: . . .

Defined in XXX.v

[u] forward Subordinate tactics: . . .

Superordinate tactics: . . .

Defined in XXX.

[u] forward_call ARGS Tactic for stepping through the call to a specified function, where *ARGS* is the product of arguments, i.e. an in-order instantiation of the function specification's WITH clause (Note: little

²or of an assertion-level entailment's LHS?

type checking is applied to *ARGS* and the tactic may behave erratically if incorrect/wrongly-typed arguments are supplied).

Always leaves a *semax*-subgoal for the program continuation. Frequently, leaves an entailment subgoal for establishing the function precondition, and/or a subgoal for a typing side condition.

Subordinate tactics: ...

Superordinate tactics: None.

Defined in XXX.v, with the underlying proof rules being defined in YYY.v

[u] forward_for Including its variants for simple bounds etc

Is expected to leave the following subgoals:...

Subordinate tactics: ...

Superordinate tactics: ...

Defined in XXX

[u] forward_seq Subordinate tactics: ...

Superordinate tactics: ...

Defined in XXX.

[i] mkConciseDelta V G F Delta Applicable to a proof state with a *semax* goal. Simplifies the Δ component to the application of a sequence of initializations to the host function's *func_tycontext*.

Superordinate tactics: *semax_subcommand*

Defined in *semax_tactics.v*.

[u] semax_subcommand V G F Applicable to a proof state with a *semax* goal. Extracts the current proof state as a stand-alone statement that can be copy-and pasted to a separate file. The three arguments are

as in the statement of surrounding semax-body lemma, i.e. are of type *varspecs*, *funspecs*, and *function*.

Subordinate tactic: *mkConciseDelta*. Defined in `semax_tactics.v`.

40 *BOUNDARY*

Everything beyond here is junk, left over from the old Verifiable C manual, that needs to be rewritten and moved above the boundary.

41 *Getting started*

This *summary reference manual* is a brief guide to the VST Separation Logic for the C language. The Verified Software Toolchain and the principles of its program logics are described in the book:

Program Logics for Certified Compilers,

by Andrew W. Appel *et al.*, Cambridge University Press, 2014.

TO INSTALL THE VST SEPARATION LOGIC FOR C LIGHT:

1. Get VST from vst.cs.princeton.edu/download, or get the bleeding-edge version from <https://github.com/PrincetonUniversity/VST>.
2. Examine `vst/compccert/VERSION` to determine which version of CompCert to download. The VST comes with a copy of the CompCert front-end, in `vst/compccert/`, but (at present) CompCert's `clightgen` utility is not buildable from just the front-end distributed with VST. You'll need `clightgen` to translate `.c` files into `.v` files containing C light abstract syntax. Thus it's recommended to download and build CompCert.
3. Get CompCert from compccert.inria.fr/download.html and run `./configure` to list configurations. Select the correct option for your machine, then run `./configure <option>` followed by `make clightgen`. Create a file `vst/CONFIGURE` containing a definition for CompCert's location; if `vst` and CompCert are installed in the same parent directly, use `COMPCERT=../compccert`
 If you have not installed CompCert, use the CompCert front-end packaged with VST. Do not create a `CONFIGURE` file, and do:
`cd vst/compccert; ./make`
4. In the `vst` directory, `make`.

See also the file `vst/BUILD_ORGANIZATION`.

The `verif_reverse.v` example is described in PLCC Chapter 27. You might find it interesting to open this in the IDE, using the command shown above, and interactively step through the definitions and proofs.

Before doing proofs of your own, you may find it helpful to step through this tutorial on C light expressions and assertions:

```
cd examples/floyd_tut; coqide tutorial.v
```

(this tutorial sets up its own load paths.)

42 Differences from PLCC

71

The book *Program Logics for Certified Compilers* (Cambridge University Press, early 2014) describes *Verifiable C* version 1.1. More recent VST versions differ in the following ways from what the PLCC book describes:

- In the LOCAL component of an assertion, `temp i v` is the recommended way to write ``(eq v) (eval_id i)`, and `var i t v` is the recommended way to write ``(eq v) (eval_var i t)`. See [Chapter 55](#) of this manual.
- The type-checker now has a more refined view of `char` and short types (see [Chapter 57](#) of this manual).
- `field_mapsto` is now called `field_at`, and it is dependently typed; see [Chapter 68](#) of this manual.
- `typed_mapsto` is renamed to `data_at`, and last two arguments are swapped.
- `umapsto` (“untyped mapsto”) no longer exists.
- `mapsto sh t v w` now permits either ($w = \text{Vundef}$) or the value w belongs to type t . This permits describing uninitialized locations, i.e., `mapsto_sh t v = mapsto_sh t v Vundef`. See [Chapter 68](#) of this manual.
- Supercanonical form is now suggested; see [Chapter 55](#) of this manual.
- For function calls, do not use `forward` (except to get advice about the witness type); instead, use `forward_call`. See [page 108](#).
- C functions may now fall through the end of the function body, and this is (per the C semantics) equivalent to a `return; statement`.

43 *Memory predicates*

The axiomatic semantics (Hoare Logic of Separation) treats memories abstractly. One never has a variable m of type *memory*. Instead, one uses the Hoare Logic to manipulate predicates P on memories. Our type of “memory predicates” is called `mpred`

Although intuitively `mpred` “feels like” the type $\text{memory} \rightarrow \text{Prop}$, the underlying semantic model is different; thus we keep the type `mpred` abstract (opaque). See *Program Logics for Certified Compilers (PLCC)* for more explanation.

On the type `mpred` we form a natural deduction system `NatDed(mpred)` with conjunction `&&`, disjunction `||`, etc.; a separation logic `SepLog(mpred)` with separating conjunction `*` and `emp`; and an indirection theory `Indir(mpred)` with `▷` “later.”

The natural deduction system has a sequent (entailment) operator written $P \multimap Q$ in Coq (written $P \vdash Q$ in print), where $P, Q : \text{mpred}$. We write b entailment simply as $P = Q$ since we assume axioms of extensionality.

44 Separation Logic

73
(see PLCC Chapter 12)

```
Class NatDed (A: Type) := mkNatDed {
  andp: A → A → A;      (Notation &&)
  orp: A → A → A;        (Notation ||)
  exp: ∀ {T:Type}, (T → A) → A;    (Notation EX)
  allp: ∀ {T:Type}, (T → A) → A;    (Notation ALL)
  imp: A → A → A;        (Notation -->, here written →)
  prop: Prop → A;        (Notation !!)
  derives: A → A → Prop;    (Notation |--, here written ⊢)
  pred_ext: ∀ P Q, P ⊢ Q → Q ⊢ P → P=Q;
  derives_refl: ∀ P, P ⊢ P;
  derives_trans: ∀ {P Q R}, P ⊢ Q → Q ⊢ R → P ⊢ R;
  TT := !!True;
  FF := !!False;
  andp_right: ∀ X P Q:A, X ⊢ P → X ⊢ Q → X ⊢ (P&&Q);
  andp_left1: ∀ P Q R:A, P ⊢ R → P&&Q ⊢ R;
  andp_left2: ∀ P Q R:A, Q ⊢ R → P&&Q ⊢ R;
  orp_left: ∀ P Q R, P ⊢ R → Q ⊢ R → P||Q ⊢ R;
  orp_right1: ∀ P Q R, P ⊢ Q → P ⊢ Q||R;
  orp_right2: ∀ P Q R, P ⊢ R → P ⊢ Q||R;
  exp_right: ∀ {B: Type}(x:B)(P:A)(Q: B→A), P ⊢ Q x → P ⊢ EX x:B, Q;
  exp_left: ∀ {B: Type}(P:B→A)(Q:A), (∀ x, P x ⊢ Q) → EX x:B, P ⊢ Q;
  allp_left: ∀ {B}(P: B → A) x Q, P x ⊢ Q → ALL x:B, P ⊢ Q;
  allp_right: ∀ {B}(P: A)(Q:B→A), (∀ v, P ⊢ Q v) → P ⊢ ALL x:B, Q;
  imp_andp_adjoint: ∀ P Q R, P&&Q ⊢ R ↔ P ⊢ (Q→R);
  prop_left: ∀ (P: Prop) Q, (P → (TT ⊢ Q)) → !!P ⊢ Q;
  prop_right: ∀ (P: Prop) Q, P → (Q ⊢ !!P);
  not_prop_right: ∀ (P:A)(Q:Prop), (Q → (P ⊢ FF)) → P ⊢ !(~Q)
}.
```

```

Class SepLog (A: Type) {ND: NatDed A} := mkSepLog {
  emp: A;
  sepcon: A → A → A;      (Notation * )
  wand: A → A → A;        (Notation -*; here written →*)
  ewand: A → A → A;        (no notation; here written →◦)
  sepcon_assoc: ∀ P Q R, (P * Q) * R = P * (Q * R);
  sepcon_comm: ∀ P Q, P * Q = Q * P;
  wand_sepcon_adjoint: ∀ (P Q R: A), P * Q ⊢ R ↔ P ⊢ Q -* R;
  sepcon_andp_prop: ∀ P Q R, P * (!Q && R) = !Q && (P * R);
  sepcon_derives: ∀ P P' Q Q' : A, P ⊢ P' → Q ⊢ Q' → P * Q ⊢ P' * Q';
  ewand_sepcon: ∀ (P Q R : A), (P * Q) →◦ R = P →◦ (Q →◦ R);
  ewand_TT_sepcon: ∀ (P Q R: A),
    (P * Q) && (R →◦ TT) ⊢ (P && (R →◦ TT)) * (Q && (R →◦ TT));
  exclude_elsewhere: ∀ P Q: A, P * Q ⊢ (P && (Q →◦ TT)) * Q;
  ewand_conflict: ∀ P Q R, P * Q ⊢ FF → P && (Q →◦ R) ⊢ FF
}.

```

```

Class Indir (A: Type) {ND: NatDed A} := mkIndir {
  later: A → A; (Notation ▷)
  now_later: ∀ P: A, P ⊢ ▷ P;
  later_K: ∀ P Q, ▷ (P → Q) ⊢ (▷ P → ▷ Q);
  later_allp: ∀ T (F: T → A), ▷ (ALL x:T, F x) = ALL x:T, ▷ (F x);
  later_exp: ∀ T (F: T → A), EX x:T, ▷ (F x) ⊢ ▷ (EX x: F x);
  later_exp': ∀ T (any:T) F, ▷ (EX x: F x) = EX x:T, ▷ (F x);
  later_imp: ∀ P Q, ▷ (P → Q) = (▷ P → ▷ Q);
  loeb: ∀ P, ▷ P ⊢ P → TT ⊢ P
}.

```

```

Class SepIndir (A: Type) {NA: NatDed A}{SA: SepLog A}{IA: Indir A} :=
  mkSepIndir {
    later_sepcon: ∀ P Q, ▷ (P * Q) = ▷ P * ▷ Q;
    later_wand: ∀ P Q, ▷ (P -* Q) = ▷ P -* ▷ Q;
    later_ewand: ∀ P Q, ▷ (P →◦ Q) = (▷ P) →◦ (▷ Q)
  }.

```

45 *Mapsto and func_ptr* (see PLCC section 24) ⁷⁵

Aside from the standard operators and axioms of separation logic, we have exactly two primitive memory predicates:

Parameter `address_mapsto`:

`memory_chunk` \rightarrow `val` \rightarrow `share` \rightarrow `share` \rightarrow `address` \rightarrow `mpred`.

Parameter `func_ptr` : `funspec` \rightarrow `val` \rightarrow `mpred`.

`func_ptr` ϕ `v` means that value v is a pointer to a function with specification ϕ .

`address_mapsto` expresses what is typically written $x \mapsto y$ in separation logic, that is, a singleton heap containing just value y at address x . But we almost always use one of the following derived forms:

`mapsto` (sh :`share`) (t :`type`) (v `w`:`val`) : `mpred` describes a singleton heap with just one value w of (C-language) type t at address v , with permission-share sh .

`mapsto_` (sh :`share`) (t :`type`) (v :`val`) : `mpred` describes an *uninitialized* singleton heap with space to hold a value of type t at address v , with permission-share sh .

`field_at` (sh : `share`) (t : `type`) (f : `list ident`) (w : `reptype` (`nested_field_type2` f)) (v : `val`) describes a heap that holds just field `fld` of struct-value v , belonging to struct-type t , containing value w . If type t describes a nested struct type, then f can actually be a path of field selections that descends into the nested structures. If f is the empty path, then the field is equivalent to `data_at`. The type of w is a dependent type. *Note: arguments w, v are swapped compared to the PLCC book.*

`field_at_` (sh : `share`) (t : `type`) (fld : `ident`) (v : `val`) : `mpred` is the corresponding uninitialized structure-field.

The CompCert verified C compiler translates standard C source programs into an abstract syntax for *CompCert C*, and then translates that into abstract syntax for *C light*. Then VST Separation Logic is applied to the C light abstract syntax. C light programs proved correct using the VST separation logic can then be compiled (by CompCert) to assembly language.

C light syntax is defined by these Coq files from CompCert:

Integers. 32-bit (and 8-bit, 16-bit, 64-bit) signed/unsigned integers.

Floats. IEEE floating point numbers.

Values. The val type: integer + float + pointer + undefined.

AST. Generic support for abstract syntax.

Ctypes. C-language types and structure-field-offset computations.

Cop. Semantics of C-language arithmetic operators.

Clight. Abstract syntax of C-light expressions, statements, and functions.

veric.expr. (from VST, not CompCert) Semantics of expression evaluation.

Some of the important types and operators are described over the next few pages.

47 Verifiable C programming ⁷⁷ See PLCC

Chapter 22

In writing Verifiable C programs you must:

- Make each dereference into a top level expression (PLCC page 143)
- Make most pointer comparisons into a top level expression (PLCC page 145)
- Remove casts between int and pointer types (result in values that crash if used)

The clightgen tool automatically:

- Factors function calls into top level expressions
- Factors logical and/or operators into if statements (to capture short circuiting behavior)

Proof automation detects these two transformations and processes them with a single tactic application.

If your program uses malloc or free, you must declare and specify these as external functions. If you don't want to keep track of the size of each allocated object, you may want to change the interface of the free function. We do this in our example definitions of malloc and free in progs/queue.c and their specifications in progs/verif_queue.v.

48 32-bit Integers

78
(compcert/lib/Integers.v)

The VST program logic uses CompCert's 32-bit integer type.

Inductive comparison := Ceq | Cne | Clt | Cle | Cgt | Cge.

Definition wordsize: nat := 32. (* also instantiations for 8, 16, 64 *)

Definition modulus : Z := two_power_nat wordsize.

Definition half_modulus : Z := modulus / 2.

Definition max_unsigned : Z := modulus - 1.

Definition max_signed : Z := half_modulus - 1.

Definition min_signed : Z := -half_modulus.

Parameter int : Type.

Parameter unsigned : int → Z.

Parameter signed : int → Z.

Parameter repr : Z → int.

Definition zero := repr 0.

Definition eq (x y: int) : bool.

Definition lt (x y: int) : bool.

Definition ltu (x y: int) : bool.

Definition neg (x: int): int := repr (- unsigned x).

Definition add (x y: int): int := repr (unsigned x + unsigned y).

Definition sub (x y: int): int := repr (unsigned x - unsigned y).

Definition mul (x y: int): int := repr (unsigned x * unsigned y).

Definition divs (x y: int) : int.

Definition mods (x y: int) : int.

Definition divu (x y: int) : int.

Definition modu (x y: int) : int.

Definition and (x y: int): int := bitwise_binop andb x y.

Definition or (x y: int): int := bitwise_binop orb x y.

Definition xor (x y: int) : int := bitwise_binop xorb x y.

Definition not (x: int) : int := xor x mone.

Definition shl (x y: int): int.

Definition shru (x y: int): int.

Definition shr ($x\ y: \text{int}$): int .

Definition rol ($x\ y: \text{int}$) : int .

Definition ror ($x\ y: \text{int}$) : int .

Definition rolm ($x\ a\ m: \text{int}$): int .

Definition cmp ($c: \text{comparison}$) ($x\ y: \text{int}$) : bool .

Definition cmpu ($c: \text{comparison}$) ($x\ y: \text{int}$) : bool .

Lemma eq_dec: $\forall (x\ y: \text{int}), \{x = y\} + \{x <> y\}$.

Theorem unsigned_range: $\forall i, 0 \leq \text{unsigned } i < \text{modulus}$.

Theorem unsigned_range_2: $\forall i, 0 \leq \text{unsigned } i \leq \text{max_unsigned}$.

Theorem signed_range: $\forall i, \text{min_signed} \leq \text{signed } i \leq \text{max_signed}$.

Theorem repr_unsigned: $\forall i, \text{repr } (\text{unsigned } i) = i$.

Lemma repr_signed: $\forall i, \text{repr } (\text{signed } i) = i$.

Theorem unsigned_repr:

$\forall z, 0 \leq z \leq \text{max_unsigned} \rightarrow \text{unsigned } (\text{repr } z) = z$.

Theorem signed_repr:

$\forall z, \text{min_signed} \leq z \leq \text{max_signed} \rightarrow \text{signed } (\text{repr } z) = z$.

Theorem signed_eq_unsigned:

$\forall x, \text{unsigned } x \leq \text{max_signed} \rightarrow \text{signed } x = \text{unsigned } x$.

Theorem unsigned_zero: $\text{unsigned zero} = 0$.

Theorem unsigned_one: $\text{unsigned one} = 1$.

Theorem signed_zero: $\text{signed zero} = 0$.

Theorem eq_sym: $\forall x\ y, \text{eq } x\ y = \text{eq } y\ x$.

Theorem eq_spec: $\forall (x\ y: \text{int}), \text{if } \text{eq } x\ y \text{ then } x = y \text{ else } x <> y$.

Theorem eq_true: $\forall x, \text{eq } x\ x = \text{true}$.

Theorem eq_false: $\forall x\ y, x <> y \rightarrow \text{eq } x\ y = \text{false}$.

Theorem add_unsigned: $\forall x\ y, \text{add } x\ y = \text{repr } (\text{unsigned } x + \text{unsigned } y)$.

Theorem add_signed: $\forall x\ y, \text{add } x\ y = \text{repr } (\text{signed } x + \text{signed } y)$.

Theorem add_commut: $\forall x\ y, \text{add } x\ y = \text{add } y\ x$.

Theorem add_zero: $\forall x, \text{add } x\ \text{zero} = x$.

Theorem add_zero_l: $\forall x, \text{add zero } x = x$.

Theorem add_assoc: $\forall x\ y\ z, \text{add } (\text{add } x\ y)\ z = \text{add } x\ (\text{add } y\ z)$.

Theorem neg_repr: $\forall z, \text{neg} (\text{repr } z) = \text{repr } (-z).$

Theorem neg_zero: $\text{neg zero} = \text{zero}.$

Theorem neg_involutive: $\forall x, \text{neg} (\text{neg } x) = x.$

Theorem neg_add_distr: $\forall x y, \text{neg}(\text{add } x y) = \text{add} (\text{neg } x) (\text{neg } y).$

Theorem sub_zero_l: $\forall x, \text{sub } x \text{ zero} = x.$

Theorem sub_zero_r: $\forall x, \text{sub zero } x = \text{neg } x.$

Theorem sub_add_opp: $\forall x y, \text{sub } x y = \text{add } x (\text{neg } y).$

Theorem sub_idem: $\forall x, \text{sub } x x = \text{zero}.$

Theorem sub_add_l: $\forall x y z, \text{sub} (\text{add } x y) z = \text{add} (\text{sub } x z) y.$

Theorem sub_add_r: $\forall x y z, \text{sub } x (\text{add } y z) = \text{add} (\text{sub } x z) (\text{neg } y).$

Theorem sub_shifted: $\forall x y z, \text{sub} (\text{add } x z) (\text{add } y z) = \text{sub } x y.$

Theorem sub_signed: $\forall x y, \text{sub } x y = \text{repr} (\text{signed } x - \text{signed } y).$

Theorem mul_commut: $\forall x y, \text{mul } x y = \text{mul } y x.$

Theorem mul_zero: $\forall x, \text{mul } x \text{ zero} = \text{zero}.$

Theorem mul_one: $\forall x, \text{mul } x \text{ one} = x.$

Theorem mul_assoc: $\forall x y z, \text{mul} (\text{mul } x y) z = \text{mul } x (\text{mul } y z).$

Theorem mul_add_distr_l: $\forall x y z, \text{mul} (\text{add } x y) z = \text{add} (\text{mul } x z) (\text{mul } y z).$

Theorem mul_signed: $\forall x y, \text{mul } x y = \text{repr} (\text{signed } x * \text{signed } y).$

and many more axioms for the bitwise operators, shift operators, signed/unsigned division and mod operators.

49 C expression syntax

(compcert/cfrontend/Clight.v)

Inductive $\text{expr} : \text{Type} :=$

$(* 1 *)$ | $\text{Econst_int}: \text{int} \rightarrow \text{type} \rightarrow \text{expr}$
 $(* 1.0 *)$ | $\text{Econst_float}: \text{float} \rightarrow \text{type} \rightarrow \text{expr}$ (** double precision **)
 $(* 1.0f0 *)$ | $\text{Econst_single}: \text{float} \rightarrow \text{type} \rightarrow \text{expr}$ (** single precision **)
 $(* 1L *)$ | $\text{Econst_long}: \text{int64} \rightarrow \text{type} \rightarrow \text{expr}$
 $(* x *)$ | $\text{Evar}: \text{ident} \rightarrow \text{type} \rightarrow \text{expr}$
 $(* x *)$ | $\text{Etempvar}: \text{ident} \rightarrow \text{type} \rightarrow \text{expr}$
 $(* *e *)$ | $\text{Ederef}: \text{expr} \rightarrow \text{type} \rightarrow \text{expr}$
 $(* \&e *)$ | $\text{Eaddrof}: \text{expr} \rightarrow \text{type} \rightarrow \text{expr}$
 $(* \sim e *)$ | $\text{Eunop}: \text{unary_operation} \rightarrow \text{expr} \rightarrow \text{type} \rightarrow \text{expr}$
 $(* e + e *)$ | $\text{Ebinop}: \text{binary_operation} \rightarrow \text{expr} \rightarrow \text{expr} \rightarrow \text{type} \rightarrow \text{expr}$
 $(* (int)e *)$ | $\text{Ecast}: \text{expr} \rightarrow \text{type} \rightarrow \text{expr}$
 $(* e.f *)$ | $\text{Efield}: \text{expr} \rightarrow \text{ident} \rightarrow \text{type} \rightarrow \text{expr}$.

Definition $\text{typeof} (e: \text{expr}) : \text{type} :=$

match e **with**

| $\text{Econst_int } _ty \Rightarrow ty$
 | $\text{Econst_float } _ty \Rightarrow ty$
 | $\text{Evar } _ty \Rightarrow ty$
 | ... *et cetera*.

```
Function bool_val (v: val) (t: type) : option bool :=
  match classify_bool t with
  | bool_case_i =>
    match v with
    | Vint n => Some (negb (Int.eq n Int.zero))
    | _ => None
    end
  | bool_case_f =>
    match v with
    | Vfloat f => Some (negb (Float.cmp Ceq f Float.zero))
    | _ => None
    end
  | bool_case_p =>
    match v with
    | Vint n => Some (negb (Int.eq n Int.zero))
    | Vptr b ofs => Some true
    | _ => None
    end
  | bool_default => None
end.
```

```
Function sem_neg (v: val) (ty: type) : option val :=
  match classify_neg ty with
  | neg_case_i sg =>
    match v with
    | Vint n => Some (Vint (Int.neg n))
    | _ => None
    end
  | neg_case_f =>
    match v with
    | Vfloat f => Some (Vfloat (Float.neg f))
    | _ => None
    end
```

```
| neg_default ⇒ None
end.
```

Function sem_add (v1:val) (t1:type) (v2: val) (t2:type) : option val :=

```
match classify_add t1 t2 with
```

```
| add_case_ii sg ⇒ (**r integer addition *)
```

```
  match v1, v2 with
```

```
  | Vint n1, Vint n2 ⇒ Some (Vint (Int.add n1 n2))
```

```
  | -, - ⇒ None
```

```
  end
```

```
| add_case_ff ⇒ (**r float addition *)
```

```
  match v1, v2 with
```

```
  | Vfloat n1, Vfloat n2 ⇒ Some (Vfloat (Float.add n1 n2))
```

```
  | -, - ⇒ None
```

```
  end
```

```
| add_case_if sg ⇒ (**r int plus float *)
```

```
  match v1, v2 with
```

```
  | Vint n1, Vfloat n2 ⇒ Some (Vfloat (Float.add (cast_int_float sg n1) n2))
```

```
  | -, - ⇒ None
```

```
  end
```

```
| ... (cases omitted)
```

```
| add_case_ip ty _ ⇒ (**r integer plus pointer *)
```

```
  match v1, v2 with
```

```
  | Vint n1, Vptr b2 ofs2 ⇒
```

```
    Some (Vptr b2 (Int.add ofs2 (Int.mul (Int.repr (sizeof ty)) n1)))
```

```
  | -, - ⇒ None
```

```
  end
```

```
| add_default ⇒ None
```

```
end.
```

Function sem_sub (v1:val) (t1:type) (v2: val) (t2:type) : option val.

Function sem_mul (v1:val) (t1:type) (v2: val) (t2:type) : option val.

Function sem_div (v1:val) (t1:type) (v2: val) (t2:type) : option val.

Function sem_mod (v1:val) (t1:type) (v2: val) (t2:type) : option val.

Function sem_and (v1:val) (t1:type) (v2: val) (t2:type) : option val.

```

Function sem_cmp (c:comparison)
    (v1: val) (t1: type) (v2: val) (t2: type)
    (m: mem): option val :=
match classify_cmp t1 t2 with
| cmp_case.ii Signed =>
    match v1,v2 with
    | Vint n1, Vint n2 => Some (Val.of_bool (Int.cmp c n1 n2))
    | -, - => None
    end
| ... (many more cases )
end.

```

```

Definition sem_binary_operation
    (op: binary_operation)
    (v1: val) (t1: type) (v2: val) (t2:type)
    (m: mem): option val :=
match op with
| Oadd => sem_add v1 t1 v2 t2
| Osub => sem_sub v1 t1 v2 t2
| Omul => sem_mul v1 t1 v2 t2
| Omod => sem_mod v1 t1 v2 t2
| Odiv => sem_div v1 t1 v2 t2
| Oand => sem_and v1 t1 v2 t2
| Oor => sem_or v1 t1 v2 t2
| Oxor => sem_xor v1 t1 v2 t2
| Oshl => sem_shl v1 t1 v2 t2
| Oshr => sem_shr v1 t1 v2 t2
| Oeq => sem_cmp Ceq v1 t1 v2 t2 m
| One => sem_cmp Cne v1 t1 v2 t2 m
| Olt => sem_cmp Clt v1 t1 v2 t2 m
| Ogt => sem_cmp Cgt v1 t1 v2 t2 m
| Ole => sem_cmp Cle v1 t1 v2 t2 m
| Oge => sem_cmp Cge v1 t1 v2 t2 m
end.

```

51 C expression evaluation

(vst/veric/expr.v)

Definition eval_id (id: ident) (ρ : environ).

(look up the temporary variable ``id'' in ρ *)*

Definition eval_cast (t t': type) (v: val) : val.

(cast value v from type t to type t', but beware! There are
be three types involved, if you include the native type of v. *)*

Definition eval_unop (op: Cop.unary_operation) (t1 : type) (v1 : val) : val.

Definition eval_binop (op: Cop.binary_operation)

(t1 t2 : type) (v1 v2: val) : val.

Definition force_ptr (v: val) : val :=

match v **with** Vptr l ofs \Rightarrow v | _ \Rightarrow Vundef **end**.

Definition eval_struct_field (delta: Z) (v: val) : val.

(offset the pointer-value v by delta *)*

Definition eval_field (ty: type) (fld: ident) (v: val) : val.

(calculate the lvalue of (but do not fetch/dereference!)
a structure/union field of value v *)*

Definition eval_var (id:ident) (ty: type) (rho: environ) : val.

(Get the lvalue (address of) an addressable local variable
(if there is one of that name) or else a global variable *)*

Definition deref_noload (ty: type) (v: val) : val.

(For By-reference types such as arrays that dereference
without actually fetching *)*

match access_mode ty **with** By_reference \Rightarrow v | _ \Rightarrow Vundef **end**.

Fixpoint eval_expr (e: expr) : environ \rightarrow val :=

match e **with**

| Econst_int i ty \Rightarrow `(Vint i)
 | Econst_float f ty \Rightarrow `(Vfloat f)
 | Etempvar id ty \Rightarrow eval_id id
 | Eaddrof a ty \Rightarrow eval_lvalue a
 | Eunop op a ty \Rightarrow `(eval_unop op (typeof a)) (eval_expr a)
 | Ebinop op a1 a2 ty \Rightarrow
 `(eval_binop op (typeof a1) (typeof a2))
 (eval_expr a1) (eval_expr a2)
 | Ecast a ty \Rightarrow `(eval_cast (typeof a) ty) (eval_expr a)
 | Evar id ty \Rightarrow `(deref_noload ty) (eval_var id ty)
 | Ederef a ty \Rightarrow `(deref_noload ty) (`force_ptr (eval_expr a))
 | Efield a i ty \Rightarrow `(deref_noload ty)
 (`(eval_field (typeof a) i) (eval_lvalue a))

end

with eval_lvalue (e: expr) : environ \rightarrow val :=

match e **with**

| Evar id ty \Rightarrow eval_var id ty
 | Ederef a ty \Rightarrow `force_ptr (eval_expr a)
 | Efield a i ty \Rightarrow `(eval_field (typeof a) i) (eval_lvalue a)
 | _ \Rightarrow `Vundef

end.

52 *C* type checking

87
(See PLCC Chapter 25)

Ideally, you will never notice the typechecker, but it may occasionally generate side conditions that can not be solved automatically. If you get a proof goal from the typechecker, it will be an entailment $P \vdash \text{denote_tc_assert } (\dots)$. PLCC Chapter 26 discusses what you can do to solve these goals.

If you are asked to prove an entailment where the typechecking condition evaluates to False, this may be because your program is not written in Verifiable C. You may need to perform some local transformations on your C program in order to proceed. We listed these transformations on [page 77](#).

The type-context will always be visible in your proof in a line that looks like `Delta := abbreviate : tycontext`. The `abbreviate` hides the implementation of the type context (which is generally large and uninteresting). The `query_context` tactic shows the result of looking up a variable in a typecontext. The tactic `query_context Delta _p.` will add hypothesis `QUERY : (temp_types Delta) ! _p = Some (tptr t_struct_list, true)`. This means that in `Delta`, `_p` is a temporary variable with type `tptr t_struct_list` and that it is known to be initialized.

53 Lifted separation logic

Chapter 21)

88
(See PLCC

Assertions in our Hoare triple of separation are presented as $\text{env} \rightarrow \text{mpred}$, that is, functions from environment to memory-predicate, using our natural deduction system $\text{NatDed}(\text{mpred})$ and separation logic $\text{SepLog}(\text{mpred})$.

Given a separation logic over a type B of formulas, and an arbitrary type A , we can define a *lifted* separation logic over functions $A \rightarrow B$. The operations are simply lifted pointwise over the elements of A . Let $P, Q : A \rightarrow B$, let $R : T \rightarrow A \rightarrow B$ then define,

$$\begin{aligned}(P \&\& Q) : A \rightarrow B &:= \text{fun } a \Rightarrow Pa \&\& Qa \\(P \parallel Q) : A \rightarrow B &:= \text{fun } a \Rightarrow Pa \parallel Qa \\(\exists x. R(x)) : A \rightarrow B &:= \text{fun } a \Rightarrow \exists x. Rxa \\(\forall x. R(x)) : A \rightarrow B &:= \text{fun } a \Rightarrow \forall x. Rxa \\(P \longrightarrow Q) : A \rightarrow B &:= \text{fun } a \Rightarrow Pa \longrightarrow Qa \\(P \vdash Q) : A \rightarrow B &:= \forall a. Pa \vdash Qa \\(P * Q) : A \rightarrow B &:= \text{fun } a \Rightarrow Pa * Qa \\(P \multimap Q) : A \rightarrow B &:= \text{fun } a \Rightarrow Pa \multimap Qa\end{aligned}$$

In Coq we formalize the typeclass instances LiftNatDed , LiftSepLog , etc., as shown below. For a type B , whenever $\text{NatDed } B$ and $\text{SepLog } B$ (and so on) have been defined, the lifted instances $\text{NatDed } (A \rightarrow B)$ and $\text{SepLog } (A \rightarrow B)$ (and so on) are automagically provided by the typeclass system.

Instance $\text{LiftNatDed}(A\ B : \text{Type})\{\text{ND} : \text{NatDed } B\} : \text{NatDed } (A \rightarrow B) :=$
 $\text{mkNatDed } (A \rightarrow B)$

```
(*andp*) (fun P Q x => andp (P x) (Q x))
(*orp*) (fun P Q x => orp (P x) (Q x))
(*exp*) (fun {T} (F : T -> A -> B) (a : A) => exp (fun x => F x a))
(*allp*) (fun {T} (F : T -> A -> B) (a : A) => allp (fun x => F x a))
(*imp*) (fun P Q x => imp (P x) (Q x))
(*prop*) (fun P x => prop P)
(*derives*) (fun P Q => forall x, derives (P x) (Q x))
```

Instance LiftSepLog (A B: Type) {NB: NatDed B}{SB: SepLog B}
 : SepLog (A → B).
 apply (mkSepLog (A → B) _ (fun ρ ⇒ emp)
 (fun P Q ρ ⇒ P ρ * Q ρ) (fun P Q ρ ⇒ P ρ -* Q ρ)).
(fill in proofs here *)*

In particular, if P and Q are functions of type $\text{environ} \rightarrow \text{mpred}$ then we can write $P * Q$, $P \&\& Q$, and so on.

Consider this assertion:

```
fun ρ ⇒ mapsto sh tint (eval_id _x ρ) (eval_id _y ρ)
      * mapsto sh tint (eval_id _u ρ) (Vint Int.zero)
```

which might appear as the precondition of a Hoare triple. It represents $(x \mapsto y) * (u \mapsto 0)$ written in informal separation logic, where x, y, u are C-language variables of integer type. Because it can be inconvenient to manipulate explicit lambda expressions and explicit environment variables ρ , we may write it in lifted form,

```
`(mapsto sh tint) (eval_id _x) (eval_id _y)
* `(mapsto sh tint) (eval_id _u) `(Vint Int.zero)
```

Each of the first two backquotes lifts a function from type $\text{val} \rightarrow \text{val} \rightarrow \text{mpred}$ to type $(\text{environ} \rightarrow \text{val}) \rightarrow (\text{environ} \rightarrow \text{val}) \rightarrow (\text{environ} \rightarrow \text{mpred})$, and the third one lifts from val to $\text{environ} \rightarrow \text{val}$.

We write a *canonical form* of an assertion as,

$$\text{PROP}(P_0; P_1; \dots, P_{l-1}) \text{LOCAL}(Q_0; Q_1; \dots, Q_{m-1}) \text{SEP}(R_0; R_1; \dots, R_{n-1})$$

The $P_i : \text{Prop}$ are Coq propositions—these are independent of the program variables and the memory. The $Q_i : \text{environ} \rightarrow \text{Prop}$ are local—they depend on program variables but not on memory. The $R_i : \text{environ} \rightarrow \text{mpred}$ are assertions of separation logic, which may depend on both program variables and memory.

The PROP/LOCAL/SEP form is defined formally as,

Definition $\text{PROPx} (P: \text{list Prop}) (Q: \text{assert}) :=$
 $\text{andp (prop (fold_right and True P)) } Q.$

Notation $"\text{'PROP' } (x; \dots; y) z" :=$
 $(\text{PROPx (cons x\%type .. (cons y\%type nil) ..) } z) \text{ (at level 10) : logic.}$

Notation $"\text{'PROP' } () z" := (\text{PROPx nil } z) \text{ (at level 10) : logic.}$

Definition $\text{LOCALx} (Q: \text{list (environ} \rightarrow \text{Prop)}) (R: \text{assert}) :=$
 $\text{andp (local (fold_right (``and) (``True) Q)) } R.$

Notation $"\text{'LOCAL' } (x; \dots; y) z" :=$
 $(\text{LOCALx (cons x\%type .. (cons y\%type nil) ..) } z) \text{ (at level 9) : logic.}$

Notation $"\text{'LOCAL' } () z" := (\text{LOCALx nil } z) \text{ (at level 9) : logic.}$

Definition $\text{SEPx} (R: \text{list assert}) : \text{assert} := \text{fold_right sepcon emp } R.$

Notation $"\text{'SEP' } (x; \dots; y)" :=$
 $(\text{SEPx (cons x\%logic .. (cons y\%logic nil) ..) }) \text{ (at level 8) : logic.}$

Notation $"\text{'SEP' } ()" := (\text{SEPx nil}) \text{ (at level 8) : logic.}$

Notation $"\text{'SEP' } ()" := (\text{SEPx nil}) \text{ (at level 8) : logic.}$

Thus, $\text{PROP}(P_0; P_1) \text{LOCAL}(Q_0; Q_1) \text{SEP}(R_0; R_1)$ is equivalent to $\text{prop } P_0 \wedge \text{prop } P_1 \&\& \text{prop } Q_0 \&\& \text{prop } Q_1 \&\& (R_0 * R_1).$

55 Supercanonical forms

A canonical form $\text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(\vec{R})$ is *supercanonical* if:

- Every element of \vec{Q} has the form $\text{temp } i \ V$ or $\text{var } i \ t \ V$, where V is a Coq expression of type `val` and i is $\beta\eta$ -equivalent to a constant (a ground term of type `ident`). The term $\text{temp } i \ V$ (of type `environ \rightarrow Prop`) is equivalent to $\text{`}(eq \ V) \ (\text{eval_id } i)$. The term $\text{var } i \ t \ V$ (of type `environ \rightarrow Prop`) is equivalent to $\text{`}(eq \ V) \ (\text{eval_var } i \ t)$.
- Every element of \vec{R} is $\text{`}(E)$ where E is a Coq expression of type `mpred`.

When assertions (preconditions of `semax`) are kept in supercanonical form, the forward tactic for symbolic execution runs *much* faster. That is,

- forward through assignment statements (including loads/stores) is up to 10 times faster for supercanonical preconditions than for ordinary (canonical) preconditions.
- Future versions of the forward tactic may *require* the precondition to be in supercanonical form.

An entailment $\text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(\vec{R}) \vdash \text{PROP}(\vec{P}')\text{LOCAL}(\vec{Q}')\text{SEP}(\vec{R}')$ is a sequent in our *lifted* separation logic; each side has type $\text{environ} \rightarrow \text{mpred}$.

By definition of the lifted entailment \vdash it means exactly,
 $\forall \rho. \text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(\vec{R})_\rho \vdash \text{PROP}(\vec{P}')\text{LOCAL}(\vec{Q}')\text{SEP}(\vec{R}')_\rho$.

There are two ways to prove such an entailment: Explicitly introduce ρ (descend into an entailment on mpred) and unfold the $\text{PROP}/\text{LOCAL}/\text{SEP}$ form; or stay in canonical form and rewrite in the lifted logic. Either way may be appropriate; this chapter describes how to descend. The `go_lower` tactic, described on this page, is rarely called directly; it is the first step of the `entailer` tactic (page 100) when applied to lifted entailments.

The tactic `go_lower` tactic does the following:

1. `intros ?rho`, as described above.
2. If the first conjunct of the left-hand-side `LOCAL`s is `tc_envIRON Δ ρ`, move it above the line; this be useful in step 6.
3. Unfold definitions for *canonical forms* (`PROPx LOCALx SEPx`), *expression evaluation* (`eval_exprlist eval_expr eval_lvalue cast_expropt eval_cast eval_binop eval_unop`), *casting* (`eval_cast classify_cast`) *type-checking* (`tc_expropt tc_expr tc_lvalue typecheck_expr typecheck_lvalue denote_tc_assert`), *function postcondition operators* (`function_body_ret_assert make_args' bind_ret get_result1 retval`), *lifting operators* (`liftx LiftEnviron Tarrow Tend lift_S lift_T lift_prod lift_last lifted lift_uncurry_open lift_curry local lift lift0 lift1 lift2 lift3`).
4. Simplify by `simpl`.
5. Rewrite by the `rewrite-hint` environment `go_lower`, which contains just a very few rules to evaluate certain environment lookups.
6. Recognize local variables.

Local variables that appear in the lifted canonical form as `(eval_id _x)` will be replaced by Coq variables `x`, provided that: (1) \vec{Q} includes a clause of the form `(tc_envIRON Δ)`, and (2) there is a hypothesis name `x _x` “above the line.” (See PLCC section 26). In addition, a typechecking hypothesis for `x` will be introduced above the line (see Chapter 57).

57 Welltypedness of variables

The typechecker ensures some invariants about the values of C-program variables: if a variable is initialized, it contains a value of its declared type.

Function parameters (accessed by Etempvar expressions) are always initialized. Nonaddressable local variables (accessed by Etempvar expressions) and address-taken local variables (accessed by Evar) may be uninitialized or initialized. Global variables (accessed by Evar) are always initialized.

The typechecker keeps track of the initialization status of local nonaddressable variables, *conservatively*: if on all paths from function entry to the current point—assuming that the conditions on if-expressions and while-expressions are uninterpreted/nondeterministic—there is an assignment to variable x , then x is known to be initialized.

The initialization status of addressable local variables is tracked in the separation logic, by assertions such as $v \mapsto _$ or $v \mapsto i$ for uninitialized and initialized variables, respectively.

Proofs using the forward tactic will typically generate proof obligations (for the user to solve) of the form,

$$\text{ENTAIL } \Delta, \text{PROP}(\vec{P}) \text{ LOCAL}(\vec{Q}) \text{ SEP}(\vec{R}) \vdash \text{PROP}(\vec{P}') \text{ LOCAL}(\vec{Q}') \text{ SEP}(\vec{R}')$$

Δ keeps track of which nonaddressable local variables are initialized; says that all references to local variables contain values of the right type; and says that all addressable locals and globals point to an appropriate block of memory.

The `go.lower` tactic (usually) deletes the assertion `tc_envirion` Δ ρ after deriving type-checking assertions of the form `tc_val` τ v for each variable v of type τ ; it puts these assertions above the line.

Definition $\text{tc_val } (\tau : \text{type}) : \text{val} \rightarrow \text{Prop} :=$
match τ **with**
 | $\text{Tint } \text{sz } \text{sg } _ \Rightarrow \text{is_int } \text{sz } \text{sg}$
 | $\text{Tlong } _ _ \Rightarrow \text{is_long}$
 | $\text{Tfloat } \text{F64 } _ \Rightarrow \text{is_float}$
 | $\text{Tfloat } \text{F32 } _ \Rightarrow \text{is_single}$
 | $\text{Tpointer } _ _ \mid \text{Tarray } _ _ _$
 | $\text{Tfunction } _ _ _ \mid \text{Tcomp_ptr } _ _ \Rightarrow \text{is_pointer_or_null}$
 | $\text{Tstruct } _ _ _ \Rightarrow \text{isptr}$
 | $\text{Tunion } _ _ _ \Rightarrow \text{isptr}$
 | $_ \Rightarrow (\text{fun } _ \Rightarrow \text{False})$
end.

Since τ is concrete, $\text{tc_val } \tau \ v$ immediately unfolds to something like,

TC0: $\text{is_int } \text{l32 } \text{Signed } (\text{Vint } i)$
 TC1: $\text{is_int } \text{l8 } \text{Unsigned } (\text{Vint } c)$
 TC2: $\text{is_int } \text{l8 } \text{Signed } (\text{Vint } d)$
 TC3: $\text{is_pointer_or_null } p$
 TC4: $\text{isptr } q$

TC0 says that i is a 32-bit signed integer; this is a tautology, so it will be automatically deleted by `go_lower`.

TC1 says that c is a 32-bit signed integer whose value is in the range of unsigned 8-bit integers (unsigned char). TC2 says that d is a 32-bit signed integer whose value is in the range of signed 8-bit integers (signed char). These hypotheses simplify to,

TC1: $0 \leq \text{Int.unsigned } c \leq \text{Byte.max_unsigned}$
 TC2: $\text{Byte.min_signed} \leq \text{Int.signed } c \leq \text{Byte.max_signed}$

58 Normalize

The `normalize` tactic performs autorewrite **with** norm and several other transformations. Many of the simplifications performed by `normalize` on entailments (whether lifted or unlifted) can be done more efficiently and systematically by `entailer`. However, on Hoare triples, `entailer` does not apply, and `normalize` is quite appropriate.

The norm rewrite-hint database uses several sets of rules.

Generic separation-logic simplifications.

$$\begin{array}{l}
 P * \text{emp} = P \qquad \text{emp} * P = P \qquad P \&\& \top = P \qquad \top \&\& P = P \\
 (EXx : A, P) * Q = EXx : A, P * Q \qquad P * (EXx : A, Q) = EXx : A, P * Q \\
 (EXx : A, P) \&\& Q = EXx : A, P \&\& Q \\
 P \&\& (EXx : A, Q) = EXx : A, P \&\& Q \qquad P * (!!Q \&\& R) = !!Q \&\& (P * R) \\
 (!!Q \&\& P) * R = !!Q \&\& (P * R) \qquad P \&\& \perp = \perp \qquad \perp \&\& P = \perp \qquad P * \perp = \perp \\
 \perp * P = \perp \qquad P \rightarrow (!!P \&\& Q = Q) \qquad P \rightarrow (!!P = \top) \qquad P \&\& P = P \\
 (EX_ : _, P) = P \qquad \text{local 'True} = \top
 \end{array}$$

Unlifting.

$$\begin{array}{l}
 'f \ \rho = f \text{ [when } f \text{ has arity 0]} \qquad 'f \ a_1 \ \rho = f \ (a_1 \ \rho) \text{ [when } f \text{ has arity 1]} \\
 'f \ a_1 \ a_2 \ \rho = f \ (a_1 \ \rho) \ (a_2 \ \rho) \text{ [when } f \text{ has arity 2, etc.]} \qquad \text{local } P \ \rho = !! (P \ \rho) \\
 (P * Q) \rho = P \rho * Q \rho \qquad (P \&\& Q) \rho = P \rho \&\& Q \rho \qquad (!!P) \rho = !!P \\
 !! (P \wedge Q) = !!P \&\& !!Q \qquad (EXx : A, P \ x) \rho = EXx : A, P \ x \ \rho \\
 '(EXx : B, P \ x) = EXx : B, '(P \ x) \qquad '(P * Q) = 'P * 'Q \\
 '(P \&\& Q) = 'P \&\& 'Q
 \end{array}$$

Pulling nonspatial propositions out of spatial ones.

$$\text{local } P \ \&\& \ !Q = !Q \ \&\& \ \text{local } P$$

$$\text{local } P \ \&\& \ (!Q \ \&\& \ R) = !Q \ \&\& \ (\text{local } P \ \&\& \ R)$$

$$(\text{local } P \ \&\& \ Q) * R = \text{local } P \ \&\& \ (Q * R)$$

$$Q * (\text{local } P \ \&\& \ R) = \text{local } P \ \&\& \ (Q * R)$$

Canonical forms.

$$\text{local } Q_1 \ \&\& \ (\text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(\vec{R})) = \text{PROP}(\vec{P})\text{LOCAL}(Q_1; \vec{Q})\text{SEP}(\vec{R})$$

$$\text{PROP}\vec{P}\text{LOCAL}\vec{Q}\text{SEP}(!P_1; \vec{R}) = \text{PROP}(P_1; \vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(\vec{R})$$

$$\text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(\text{local } Q_1; \vec{R}) = \text{PROP}(\vec{P})\text{LOCAL}(Q_1; \vec{Q})\text{SEP}(\vec{R})$$

Modular Integer arithmetic.

$$\text{Int.sub } x \ x = \text{Int.zero}$$

$$\text{Int.sub } x \ \text{Int.zero} = x$$

$$\text{Int.add } x \ (\text{Int.neg } x) = \text{Int.zero}$$

$$\text{Int.add } x \ \text{Int.zero} = x$$

$$\text{Int.add } \text{Int.zero } x = x$$

$$x \neq y \rightarrow \text{offset_val}(\text{offset_val } v \ i) \ j = \text{offset_val } v \ (\text{Int.add } i \ j)$$

$$\text{Int.add}(\text{Int.repr } i)(\text{Int.repr } j) = \text{Int.repr}(i + j)$$

$$\text{Int.add}(\text{Int.add } z \ (\text{Int.repr } i)) \ (\text{Int.repr } j) = \text{Int.add } z \ (\text{Int.repr}(i + j))$$

$$z > 0 \rightarrow (\text{align } 0 \ z = 0)$$

$$\text{force_int}(\text{Vint } i) = i$$

Type checking and miscellaneous.

$$\text{tc_formals}((i, t) :: r) = \text{'and } ((\text{tc_val } t) (\text{eval_id } i) (\text{tc_formals } r))$$

$$\text{tc_formals nil} = \text{'T} \quad \text{tc_andp tc_TT } e = e \quad \text{tc_andp } e \text{ tc_TT} = e$$

$$\text{eval_id } x (\text{env_set } \rho \ x \ v) = v$$

$$x \neq y \rightarrow (\text{eval_id } x (\text{env_set } \rho \ y \ v) = \text{eval_id } x \ v)$$

$$\text{isptr } v \rightarrow (\text{eval_cast_neutral } v = v)$$

$$(\exists t. \text{tc_val } t \ v \wedge \text{is_pointer_type } t) \rightarrow (\text{eval_cast_neutral } v = v)$$
Expression evaluation. (autorewrite with eval, but in fact these are usually handled just by simpl or unfold.)

$$\text{deref_noload}(\text{tarray } t \ n) = (\text{fun } v \Rightarrow v) \quad \text{eval_expr}(\text{Etempvar } i \ t) = \text{eval_id } i$$

$$\text{eval_expr}(\text{Econst_int } i \ t) = \text{'(Vint } i)$$

$$\text{eval_expr}(\text{Ebinop } op \ a \ b \ t) =$$

$$\text{'(eval_binop } op \ (\text{typeof } a) \ (\text{typeof } b)) (\text{eval_expr } a) (\text{eval_expr } b)$$

$$\text{eval_expr}(\text{Eunop } op \ a \ t) = \text{'(eval_unop } op \ (\text{typeof } a)) (\text{eval_expr } a)$$

$$\text{eval_expr}(\text{Ecast } e \ t) = \text{'(eval_cast}(\text{typeof } e) \ t) (\text{eval_expr } e)$$

$$\text{eval_lvalue}(\text{Ederef } e \ t) = \text{'force_ptr } (\text{eval_expr } e)$$
Structure fields.

$$\text{field_mapsto } sh \ t \ fld \ (\text{force_ptr } v) = \text{field_mapsto } sh \ t \ fld \ v$$

$$\text{field_mapsto_ } sh \ t \ fld \ (\text{force_ptr } v) = \text{field_mapsto_ } sh \ t \ fld \ v$$

$$\text{field_mapsto } sh \ t \ x \ (\text{offset_val } v \ \text{Int.zero}) = \text{field_mapsto } sh \ t \ x \ v$$

$$\text{field_mapsto_ } sh \ t \ x \ (\text{offset_val } v \ \text{Int.zero}) = \text{field_mapsto_ } sh \ t \ x \ v$$

$$\text{memory_block } sh \ \text{Int.zero } (\text{Vptr } b \ z) = \text{emp}$$

Postconditions. (autorewrite **with** `ret_assert`.)

$$\text{normal_ret_assert } \perp \text{ ek } vl = \perp$$

$$\text{frame_ret_assert}(\text{normal_ret_assert } P) Q = \text{normal_ret_assert } (P * Q)$$

$$\text{frame_ret_assert } P \text{ emp} = P$$

$$\text{frame_ret_assert } P Q \text{ EK_return } vl = P \text{ EK_return } vl * Q$$

$$\text{frame_ret_assert}(\text{loop1_ret_assert } P Q) R =$$

$$\text{loop1_ret_assert } (P * R)(\text{frame_ret_assert } Q R)$$

$$\text{frame_ret_assert}(\text{loop2_ret_assert } P Q) R =$$

$$\text{loop2_ret_assert } (P * R)(\text{frame_ret_assert } Q R)$$

$$\text{overridePost } P (\text{normal_ret_assert } Q) = \text{normal_ret_assert } P$$

$$\text{normal_ret_assert } P \text{ ek } vl = (!!(\text{ek} = \text{EK_normal}) \&\& (!!(vl = \text{None}) \&\& P))$$

$$\text{loop1_ret_assert } P Q \text{ EK_normal } \text{None} = P$$

$$\text{overridePost } P R \text{ EK_normal } \text{None} = P$$

$$\text{overridePost } P R \text{ EK_return} = R \text{ EK_return}$$

$$\text{function_body_ret_assert } t P \text{ EK_return } vl = \text{bind_ret } vl t P$$

Function return values.

$$\text{bind_ret } (\text{Some } v) \ t \ Q = (!\text{tc_val } t \ v \ \&\& 'Q(\text{make_args}(\text{ret_temp} :: \text{nil}) \ (v :: \text{nil})))$$

$$\text{make_args}' \ \sigma \ a \ \rho = \text{make_args } (\text{map } \text{fst } (\text{fst } \sigma)) \ (a \ \rho) \ \rho$$

$$\text{make_args}(i :: l)(v :: r)\rho = \text{env_set}(\text{make_args}(l)(r)\rho) \ i \ v$$

$$\text{make_args } \text{nil } \text{nil} = \text{globals_only} \quad \text{get_result}(\text{Some } x) = \text{get_result1}(x)$$

$$\text{retval}(\text{get_result1 } i \ \rho) = \text{eval_id } i \ \rho \quad \text{retval}(\text{env_set } \rho \ \text{ret_temp } v) = v$$

$$\text{retval}(\text{make_args}(\text{ret_temp} :: \text{nil}) \ (v :: \text{nil}) \ \rho) = v$$

$$\text{ret_type}(\text{initialized } i \ \Delta) = \text{ret_type}(\Delta)$$

IN ADDITION TO REWRITING, the `normalize` tactic applies the following rules:

$$P \vdash \top \quad \perp \vdash P \quad P \vdash P * \top \quad (\forall x. (P \vdash Q)) \rightarrow (EXx : A, P \vdash Q)$$

$$(P \rightarrow (\top \vdash Q)) \rightarrow (!P \vdash Q) \quad (P \rightarrow (Q \vdash R)) \rightarrow (!P \ \&\& \ Q \vdash R)$$

and does some rewriting and substitution when P is an equality in the goal, $(P \rightarrow (Q \vdash R))$.

Given the goal $x \rightarrow P$, where x is not a Prop, the `normalize` avoids doing an intro. This allows the user to choose an appropriate name for x .

Our entailer tactic is a partial solver for entailments in the separation logic over `mpred`. If it cannot solve the goal entirely, it leaves a simplified subgoal for the user to prove. The algorithm is this:

1. Apply `go_lower` if the goal is in the lifted separation logic.
2. Gather all the pure propositions to a single pure proposition (in each of the hypothesis and conclusion).
3. Given the resulting goal $!!(P_1 \wedge \dots \wedge P_n) \&\& (Q_1 * \dots * Q_m) \vdash !!(P'_1 \wedge \dots \wedge P'_{n'}) \&\& (Q'_1 * \dots * Q'_{m'})$, move each of the pure propositions P_i “above the line.” Any P_i that’s an easy consequence of other above-the-line hypotheses is deleted. Certain kinds of P_i are simplified in some ways.
4. For each of the Q_i , `saturate_local` extracts any pure propositions that are consequences of spatial facts, and inserts them above the line if they are not already present. For example, $p \mapsto_{\tau} q$ has two pure consequences: `isptr p` (meaning that p is a pointer value, not an integer or float) and `tc_val τ q` (that the value q has type τ).
5. For any equations $(x = \dots)$ or $(\dots = x)$ above the line, substitute x .
6. Simplify C-language comparisons.
7. Rewriting: the `normalize` tactic, as explained in Chapter 14.
8. Repeat from step 2, as long as progress is made.
9. Now the proof goal has the form $(Q_1 \dots * Q_m) \vdash !!(P'_1 \wedge \dots \wedge P'_{n'}) \&\& (Q'_1 \dots * Q'_{m'})$. Any of the P'_i provable by auto are removed. If $Q_1 * \dots * Q_m \vdash Q'_1 * \dots * Q'_{m'}$ is trivially proved, then the entire $\&\& Q'_1 * \dots * Q'_{m'}$ is removed.

AT THIS POINT the entailment may have been solved entirely. Or there may be some remaining P'_i and/or Q'_i proof goals on the right hand side.

In the judgment $\Delta \vdash \{P\} c \{R\}$, written in Coq as

`semax (Δ : tycontext) (P : environ \rightarrow mpred) (c : statement) (R : ret_assert)`

Δ is a *type context*, giving types of function parameters, local variables, and global variables; and giving *specifications* (funspec) of global functions.

P is the precondition;

c is a command in the C language; and

R is the postcondition. Because a c statement can exit in different ways (fall-through, continue, break, return), a `ret_assert` has predicates for all of these cases.

The *basic* VST separation logic is specified in `vst/veric/SeparationLogic.v`, and contains rules such as,

$$\text{semax_set_forward} \frac{}{\Delta \vdash \{\triangleright P\} \ x := e \ \{\exists v. x = (e[v/x]) \wedge P[v/x]\}}$$

Axiom `semax_set_forward`: $\forall \Delta \ P \ (x: \text{ident}) \ (e: \text{expr})$,
`semax $\Delta \ (\triangleright \ (\text{local} \ (\text{tc_expr} \ \Delta \ e) \ \&\& \ \text{local} \ (\text{tc_temp_id} \ \text{id} \ (\text{typeof} \ e) \ \Delta \ e) \ \&\& \ P))$`
`(Sset $x \ e$)`
`(normal_ret_assert`
`(EX old:val,`
`local ($\text{'eq} \ (\text{eval_id} \ x) \ (\text{subst} \ x \ (\text{'old}) \ (\text{eval_expr} \ e))) \ \&\&$`
`subst $x \ (\text{'old}) \ P$)).`

However, most C-program verifications will not use the *basic* rules, but will use derived rules whose preconditions are in canonical (PROP/LOCAL/SEP) form. Furthermore, program verifications do not even use the derived rules directly, but use *symbolic execution tactics* that choose which derived rules to apply. So we will not show the rules here; we describe how to use the tactical system.

Many of the Hoare rules, such as the one on [page 101](#),

$$\text{semax_set_forward} \frac{}{\Delta \vdash \{\triangleright P\} \ x := e \ \{\exists v. x = (e[v/x]) \wedge P[v/x]\}}$$

have the operator \triangleright (pronounced “later”) in their precondition.

The modal assertion $\triangleright P$ is a slightly weaker version of the assertion P . It is used for reasoning by induction over how many steps left we intend to run the program. The most important thing to know about \triangleright later is that P is stronger than $\triangleright P$, that is, $P \vdash \triangleright P$; and that operators such as $*$, $\&\&$, ALL (and so on) commute with later: $\triangleright(P * Q) = (\triangleright P) * (\triangleright Q)$.

This means that if we are trying to apply a rule such as `semax_set_forward`; and if we have a precondition such as

`local (tc_expr Δ e) $\&\&$ \triangleright local (tc_temp_id id t Δ e) $\&\&$ ($P_1 * \triangleright P_2$)`

then we can use the rule of consequence to *weaken* this precondition to

`\triangleright (local (tc_expr Δ e) $\&\&$ local (tc_temp_id id t Δ e) $\&\&$ ($P_1 * P_2$))`

and then apply `semax_set_forward`. We do the same for many other kinds of command rules.

This weakening of the precondition is done automatically by the forward tactic, as long as there is only one \triangleright later in a row at any point among the various conjuncts of the precondition.

A more sophisticated understanding of \triangleright is needed to build proof rules for recursive data types and for some kinds of object-oriented programming; see PLCC Chapter 19.

62 Specifying a function

Chapter 27)

103
(See PLCC

Let F be a C-language function, $t_{\text{ret}} F (t_1 x_1, t_2 x_2, \dots t_n x_n) \{ \dots \}$. The formal parameters are $\vec{x} : \vec{\tau}$ (that is, $x_1 : t_1, x_2 : t_2, \dots x_n : t_n$) and the return type is t_{ret} .

Specify F with precondition $P(\vec{a} : \vec{\tau})(\vec{x} : \vec{\tau})$ and postcondition $Q(\vec{a} : \vec{\tau})(\text{retval})$ where \vec{a} are logical variables that both the precondition and the postcondition can refer to.

The x_i are C-language variable identifiers, and the t_i are C-language types (tint, tfloat, tptr(tint), etc.). The a_i are Coq variables and the τ_i are Coq types.

Definition $F_{\text{spec}} :=$

```
DECLARE _F
  WITH  $a_1 : \tau_1, \dots a_k : \tau_k$ 
  PRE [  $x_1$  OF  $t_1, \dots, x_n$  OF  $t_n$  ]  $P$ 
  POST [  $t_{\text{ret}}$  ]  $Q$ .
```

Example: for a C function, int sumlist (struct list *p);

Definition sumlist_spec :=

```
DECLARE _sumlist
  WITH sh : share, contents : list int, p: val,
  PRE [ _p OF (tptr t_struct_list)]
    local ( `(eq p) (eval_id _p))
    && `(lseg LS sh contents p nullval)
  POST [ tint ]
    local ( `(eq (Vint (sum_int contents))) retval)
    && `(lseg LS sh contents p nullval).
```

The specification itself is an object of type ident*funspec, and in some cases it can be useful to define the components separately:

Definition sumlist_funspec : funspec :=

```

WITH sh : share, contents : list int, p: val,
PRE [ _p OF (tptr t_struct_list)]
    local `(eq p) (eval_id _p))
    && `(lseg LS sh contents p nullval)
POST [ tint ]
    local `(eq (Vint (sum_int contents))) retval)
    && `(lseg LS sh contents p nullval).

```

Definition `sumlist_spec` : `ident*funspec` :=
`DECLARE _sumlist sumlist_funspec.`

The precondition may be written in *simple form*, as shown above, or in *canonical form*:

```

Definition sumlist_spec :=
DECLARE _sumlist
  WITH sh : share, contents : list int, p: val,
  PRE [ _p OF (tptr t_struct_list)]
      PROP() LOCAL `(eq p) (eval_id _p))
      SEP `(lseg LS sh contents p nullval))
  POST [ tint ]
      local `(eq (Vint (sum_int contents))) retval)
      && `(lseg LS sh contents p nullval).

```

At present, postconditions may not use PROP/LOCAL/SEP form.

63 Specifying all functions PLCC Chapter 27)

We give each function a *specification*, typically using the DECLARE/WITH/PRE/POST notation. Then we combine these together into a *global specification*:

$$\Gamma : \text{list} (\text{ident} * \text{funspec}) := (\iota_1, \phi_1) :: (\iota_2, \phi_2) :: (\iota_3, \phi_3) :: (\iota_4, \phi_4) :: \text{nil}.$$

We also make a *global variables type specification*, listing the types of all extern global variables:

$$V : \text{list} (\text{ident} * \text{type}) := (x_1, t_1) :: (x_2, t_2) :: \text{nil}$$

The *initialization values* of extern globals are not part of V , as (generally) they are not invariant over program execution—global variables can be updated by storing into them. Initializers are accessible in the precondition to the `_main` function.

C-language functions can call each other, and themselves, and access global variables. Correctness proofs of individual functions can take advantage of the specifications of all global functions and types of global variables. Thus we construct Γ and V before proving correctness of any functions.

The next step (in a program proof) is to prove correctness of each function. For each function F in a C program, CompCert clightgen produces $_F : \text{ident}.$ $\text{f_}F : \text{function}.$ where function is a record telling the parameters and locals (and their types) and the function body. The predicate `semax_body` states that F meets its specification; for each F we must prove:

Lemma `body_` $_F$: `semax_body` $V \ \Gamma \ \text{f_}F \ F_spec.$

64 Proving a function

106
(See PLCC Chapter 27)

Lemma $\text{body_}F$: $\text{semax_body } V \ \Gamma \ f_F \ F_spec.$

Proof.

`start_function.`

`name x _x.`

`name y _y.`

`name z _z.`

Then, for each function parameter and nonaddressable local variable (scalar local variable whose address is never taken), we write a name declaration; in each case, `_x` is the identifier definition that `clightgen` has created from the source-language name, and `x` is the Coq name that we wish to use for the *value* of variable `_x` at various points. The only purpose of the name tactic is to assist the `go_lower` tactic in choosing nice names.

At this point the proof goal will be a judgment of the form,

$$\text{semax} \ \Delta \ (\text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(\vec{R})) \ c \ Post.$$

We prove such judgments as follows:

1. Manipulate the precondition $\text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(\vec{R})$ until it takes a form suitable for forward symbolic execution through the first statement in the command c . (In this we are effectively using the rule of consequence.)
2. Apply a forward tactic to step into c . This will produce zero or more entailments $A \vdash B$ to prove, where A is in canonical form; and zero or more `semax` judgments to prove.
3. Prove the entailments, typically using `go_lower`; prove the judgment, i.e., back to step 1.

Each kind of `C` command has different requirements on the form of the precondition, for the forward tactic to succeed. In each of the following cases, the expression E must not contain loads, stores, side effects, function calls, or pointer comparisons. The variable x must be a nonaddressable local variable.

- $c_1; c_2$ Sequencing of two commands. The forward tactic will work on c_1 first.
- $(c_1; c_2) c_3$ In this case, forward will re-associate the commands using the seq_assoc axiom, and work on $c_1; (c_2; c_3)$.
- $x=E$; Assignment statement. Expression E must not contain memory dereferences (loads or stores using *prefix, suffix[], or -> operators). Expression E must not contain pointer-comparisons. No restrictions on the form of the precondition (except that it must be in canonical form). The expression &p→ next does not actually load or store (it just computes an address) and is permitted.
- $x= *E$; Memory load. The SEP component of the precondition must contain an item of the form $\text{(mapsto sh t) } e \ v$, where e is equivalent to $(\text{eval_expr } E)$. For example, if E is just an identifier $(\text{Etempvar } _y \ t)$, then e could be either $(\text{eval_expr } (\text{Etempvar } _y \ t))$ or $(\text{eval_id } _y)$.
- $x= a[E]$; Array load. This is just a memory load, equivalent to $x= *(a+E)$.
- $x= E \rightarrow fld$; Field load. This is equivalent to $x= *(E.fld)$ and can actually be handled by the “memory load” case, but a special-purpose field-load rule is easier to use (and will be automatically applied by the forward tactic). In this case the SEP component of the precondition must contain $\text{(field_at sh t fld) } v \ e$, where t is the structure type to which the field fld belongs, and e is equivalent to $(\text{eval_expr } E)$.
- $*E_1 = E_2$; Memory store. The SEP component of the precondition must contain an item of the form $\text{(mapsto sh t) } e_1 \ v$ or an item $\text{(mapsto_sh t) } e_1$, where e_1 is equivalent to $(\text{eval_expr } E_1)$.
- $a[E_1]=E_2$; Array store. This is equivalent to $*(a+E_1)=E_2$; and is handled by the previous case.
- $E_1 \rightarrow fld = E_2$; Field store. This can be handled by the general store case, but a special-purpose field-store rule is easier to use. The SEP component of the precondition must contain either $\text{(field_at sh t fld) } v \ e_1$ or $\text{(field_mapsto_sh t fld) } e_1$, where t is the structure type to which the field fld belongs, and e_1 is equivalent to $(\text{eval_expr } E_1)$. The share sh must be strong enough to grant write permission, that is, $\text{writable_share}(sh)$.

$x = E_1 \text{ op } E_2$; If E_1 or E_2 evaluate to *pointers*, and op is a comparison operator ($=$, \neq , $<$, \leq , $>$, \geq), then $E_1 \text{ op } E_2$ must not occur except in this special-case assignment rule. When E_1 and E_2 both have numeric values, the ordinary *assignment statement* rule applies.

Pointer comparisons are tricky in CompCert C for reasons explained at PLCC page 249; the program logic uses the `semax_ptr.compare` rule (PLCC page 164). After applying the forward tactic, the user will be left with some proof obligations: Prove that both E_1 and E_2 evaluate to allocated locations (i.e., that the precondition implies $E_1 \xrightarrow{sh_1} _ * TT$ and also implies $E_2 \xrightarrow{sh_2} _ * TT$, for any sh_1 and sh_2). If the comparison is any of $>$, $<$, \geq , \leq , prove that E_1 and E_2 both point within the same allocated object. These are preconditions for even being permitted to test the pointers for equality (or inequality). See also [page 112](#).

if (E) C_1 else C_2 No restrictions on the form of the precondition. forward will create 3 subgoals: (1) prove that the precondition entails $\text{tc_expr} \Delta E$. For many expressions E , the condition $\text{tc_expr} \Delta E$ is simply \top , which is trivial to prove. (2) the **then** clause... (3) the **else** clause... .

while (E) C For a while-loop, use the `forward_while` tactic (page ??).

return E ; No special precondition, except that the presence/absence of E must match the nonvoid/void return type of the function. The proof goal left by forward is to show that the precondition (with appropriate substitution for the abstract variable `ret_var`) entails the function's postcondition.

$x = f(a_1, \dots, a_n)$; For a function call, use `forward_call(W)`, where W is a witness, a tuple corresponding (componentwise) to the **WITH** clause of the function specification. (If you do just forward, you'll get a message with advice about the *type* of W .)

This results a proof goal to show that the precondition implies the function precondition and includes an uninstantiated variable: The Frame represents the part of the spacial precondition that is unchanged by the function call. It will generally be instantiated by a call to `cancel`.

65 Manipulating preconditions

In some cases you cannot go forward until the precondition has a certain form. For example, in ordinary separation logic we might have $\{p \neq q \wedge p \rightsquigarrow q\} x := p \rightarrow \text{tail}\{Post\}$. In order to use the proof rule for load, we must use the rule of consequence, to prove,

$$p \neq q \wedge p \rightsquigarrow q \vdash p \neq q \wedge \exists h, t. p \mapsto (h, t) * t \rightsquigarrow q$$

then instantiate the existentials; this finally gives us

$$\{p \neq q \wedge p \mapsto (h, t) * t \rightsquigarrow q\} x := p \rightarrow \text{tail}\{Post\}$$

which is provable by the standard load rule of separation logic.

Faced with the proof goal $\text{semax } \Delta \text{ (PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(\vec{R})) \text{ } c \text{ } Post$ where $\text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(\vec{R})$ does not match the requirements for forward symbolic execution, you have several choices:

- Use the rule of consequence explicitly:
apply `semax_pre` **with** $\text{PROP}(\vec{P}')\text{LOCAL}(\vec{Q}')\text{SEP}(\vec{R}')$,
then prove $\vec{P}; \vec{Q}; \vec{R} \vdash \vec{P}'; \vec{Q}'; \vec{R}'$ using `go_lower` ([page 92](#)).
- Use the rule of consequence implicitly, by using tactics that modify the precondition (and may leave entailments for you to prove).
- Do rewriting in the precondition, either directly by the standard `rewrite` and `change` tactics, or by `normalize`.
- Extract propositions and existentials from the precondition, by using `normalize` (or by applying the rules `extract_exists_pre` and `semax_extract_PROP`).

TACTICS FOR MANIPULATING PRECONDITIONS. In many of these tactics we select specific conjuncts from the SEP items, that is, the semicolon-separated list of separating conjuncts. These tactic refer to the list by zero-based position number, 0,1,2,... For example, suppose the goal is a

semax or entailment containing $\text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(a;b;c;d;e;f;g;h;i;j)$.
Then:

focus_SEP $i\ j\ k$. Bring items $\#i, j, k$ to the front of the SEP list.

focus_sep 5. *results in* $\text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(f;a;b;c;d;e;g;h;i;j)$.

focus_sep 0. *results in* $\text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(a;b;c;d;e;f;g;h;i;j)$.

focus_SEP 1 3. *results in* $\text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(b;d;a;c;e;f;g;h;i;j)$

focus_SEP 3 1. *results in* $\text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(d;b;a;c;e;f;g;h;i;j)$

gather_SEP $i\ j\ k$. Bring items $\#i, j, k$ to the front of the SEP list and conjoin them into a single element.

gather_sep 5. *results in* $\text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(f;a;b;c;d;e;g;h;i;j)$.

gather_SEP 1 3. *results in* $\text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(b*d;a;c;e;f;g;h;i;j)$

gather_SEP 3 1. *results in* $\text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(d*b;a;c;e;f;g;h;i;j)$

replace_SEP $i\ R$. Replace the i th element the SEP list with the assertion R , and leave a subgoal to prove.

replace_sep 3 R . *results in* $\text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(a;b;c;R;e;f;g;h;i;j)$.

with subgoal $\text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(d) \vdash R$.

replace_in_pre $S\ S'$. Replace S with S' anywhere it occurs in the precondition then leave $(\vec{P};\vec{Q};\vec{R}) \vdash (\vec{P};\vec{Q};\vec{R})[S'/S]$ as a subgoal.

frame_SEP $i\ j\ k$. Apply the frame rule, keeping only elements i, j, k of the SEP list. See Chapter 66.

66 The Frame rule

Separation Logic supports the Frame rule,

$$\text{Frame} \frac{\{P\} c \{Q\}}{\{P * F\} c \{Q * F\}}$$

To use this in a forward proof, suppose you have the proof goal,

$$\text{semax} \Delta \text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(R_0;R_1;R_2) \ c_1;c_2;c_3 \ \text{Post}$$

and suppose you want to “frame out” R_2 for the duration of $c_1;c_2$, and have it back again for c_3 . First you rewrite by `seq_assoc` to yield the goal

$$\text{semax} \Delta \text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(R_0;R_1;R_2) \ (c_1;c_2);c_3 \ \text{Post}$$

Then eapply `semax_seq'` to peel off the first command $(c_1;c_2)$ in the new sequence:

$$\text{semax} \Delta \text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(R_0;R_1;R_2) \ c_1;c_2 \ ?88$$

$$\text{semax} \Delta' \ ?88 \ c_3 \ \text{Post}$$

Then `frame_SEP 0 2` to retain only $R_0;R_2$.

$$\text{semax} \Delta \text{PROP}(\vec{P})\text{LOCAL}(\vec{Q})\text{SEP}(R_0;R_2) \ c_1;c_2 \ \dots$$

Now you'll see that (in the precondition of the second subgoal) the unification variable `?88` has been instantiated in such a way that R_2 is added back in.

Pointer comparisons can be split into two cases:

1. Comparisons between two expressions that evaluate to be pointers. In this case, both of the pointers must be to *allocated* objects, or the expression will not evaluate
2. Comparisons between an expression that evaluates to the null pointer and any expression that evaluates to a value with a pointer type. This expression will always evaluate

If you are sure that your pointer comparison falls into the first case, you may treat it exactly like any other expression. The proof may eventually generate a side-condition asking you to prove that one of the expressions evaluates to the null pointer. If your pointer comparison might be between two pointers, however, the expression should be factored into its own statement (PLCC page 145).

When you use forward on a pointer comparison you might get a side condition with a disjunction. The left and right sides of the disjunction correspond to the first and second type of comparison above. In simple cases, the tactic can solve the disjunction automatically.

68 Structured data

The C programming language has struct and array to represent structured data. The *Verifiable C* logic provides operators `field_at`, `array_at`, and `data_at` to describe assertions about structs and arrays.

Given a struct definition, `struct list {int head; struct list *tail;};` the `clightgen` utility produces the type `t_struct_list` describing fields `head` and `tail`. Then these assertions are all equivalent:

```
mapsto sh tint p h *
  mapsto sh (tptr t_struct_list) (offset_val p (Vint (Int.repr 4))) t

field_at sh t_struct_list [_head] h p * field_at sh t_struct_list [_tail] t p

data_at sh t_struct_list (h,t) p

field_at sh t_struct_list nil (h,t) p
```

The version using `mapsto` is correct (assuming a 32-bit configuration of CompCert) but rather ugly; the second version is useful when you want to “frame out” a particular field; the third version describes the contents of all structure-fields at once.

The `data_at` predicate is dependently typed; the *type* of its third argument (h, t) depends on the *value* of its second argument. The dependent type is expressed by the function, `retype: type → Type` that converts C-language types into Coq Types.

Here, `retype t_struct_list = val*val`, so the type of (h, t) is $(val*val)$. The value h may be `Vint(i)` or `Vundef`, and t may be `Vpointer b z`, `Vint Int.zero`, or `Vundef`. The `Vundef` values represent uninitialized data fields.

When τ is a struct type and n is a nat, the tactic `unfold_data_at n` unfolds the n th occurrence of `data_at sh τ` to a series of `field_at sh τ ($f::nil$)`, where the f are the various fields of the struct τ . For example, it would unfold

the third assertion above to look like the second one.

The forward tactic, when the next command is a load or store command, can operate directly on `data.at` assertions; it is not necessary to unfold them to individual `field.at` conjuncts. This is a new feature of VST 1.5.

WHY ARE THE ARGUMENTS BACKWARDS? We write

`field.at sh t_struct_list [head] h p` where Reynolds would have written $p.\text{head} \mapsto h$, and we write `data.at sh t_struct_list (h,t) p` where Reynolds would have written $p \hookrightarrow (h,t)$. Putting the *contents* argument before the *pointer* argument makes it easier to express identities in our lifted separation logic. That is, we commonly have formulas such as

```
^(data.at sh t_struct_list (h,t)) (eval_id _p )tptr t_struct_list))
```

which simplify to

```
^(field.at sh t_struct_list [head] h * field.at sh t_struct_list [_tail] t)
  (eval_id _p (tptr t_struct_list))
```

Expressing these equivalences with the arguments in the other order would lead to extra lambdas, which are (ironically) no fun at all.

PARTIALLY INITIALIZED DATA STRUCTURES. Consider the program

```
struct list *f(void) {
    struct list *p = (struct list *)malloc(sizeof(struct list));
    /* 1 */ p->head = 3;
    /* 2 */ p->tail = NULL;
    /* 3 */ return p;
}
```

We do not want to assume that `malloc` returns initialized memory, so at point 1 the contents of `head` and `tail` are `Vundef`. We can write this as any of the following:

```
field.at Tsh t_struct_list [head] Vundef p
    * field.at Tsh t_struct_list [_tail] Vundef p
```

```

field_at_ Tsh t_struct_list [head] p * field_at_ Tsh t_struct_list [_tail] p
data_at_ Tsh t_struct_list (Vundef,Vundef) p
data_at_ Tsh t_struct_list p

```

If malloc returns fields that—operationally—contain defined values instead of Vundef, these assertions are still valid, as they ignore the contents of the fields.

At point 2, all the assertions above are still true, but they are weaker than the “appropriate” assertion, which may be written as any of,

```

field_at_ Tsh t_struct_list [head] (Vint(Int.repr 3)) p
      * field_at_ Tsh t_struct_list [_tail] Vundef p
field_at_ Tsh t_struct_list [head] (Vint(Int.repr 3))
      * field_at_ Tsh t_struct_list [_tail] p
data_at_ Tsh t_struct_list (Vint(Int.repr 3), Vundef) p

```

At point 3, we can write either of,

```

field_at_ Tsh t_struct_list [head] (Vint(Int.repr 3)) p
      * field_at_ Tsh t_struct_list [_tail] (Vint Int.zero) p
data_at_ Tsh t_struct_list (Vint(Int.repr 3), Vint Int.zero) p

```

FULLY INITIALIZED DATA STRUCTURES. In a function precondition it is sometimes convenient to write,

```

WITH data: retype t_struct_list
PRE [_p OF tptr t_struct_list] (**)
  `(data_at Tsh t_struct_list data) (eval_id _p (tptr t_struct_list))
POST [ ... ] ...

```

If $p \rightarrow \text{head}$ and $p \rightarrow \text{tail}$ may be uninitialized, this is fine. But if the structure is known to be initialized, the precondition as written does not express this fact. One would need to add the conjunct $!(\text{is_int } (\text{fst data}) \wedge \text{is_pointer_or_null } (\text{snd data}))$ at the point marked (**).

The function `reptype' : type → Type` expresses the type of *initialized* data structures. For example, `reptype' t_struct_list` is `(int*val)`. The function

`repinj (t: type): reptype' t → reptype t`

expresses injections from (possibly) undefined to defined values. Suppose `(h: int, t: val)` is a value of type `reptype' t_struct_list`. Then

`repinj t_struct_list (h,t) = (Vint h, t)`

Using `reptype'` one could write,

```
WITH data: reptype' t_struct_list
PRE [_p OF tptr t_struct_list]
  !(is_pointer_or_null (snd data) &&
    `(data_at Tsh t_struct_list (repinj _data)) (eval_id _p (tptr t_struct_list)))
POST [ ... ] ...
```

Notice that this only solves half the problem—for integers but not for pointers. Since defined pointers can be either NULL or a Vpointer, we use `val` to represent them, and the Coq type alone does not express the refinement. One could imagine a version of `reptype'` that uses a refinement type to accomplish this, but it might be unwieldy.

69 For loops

The C-language for loop has the general form,

for (*init*; *test*; *incr*) *body*

To solve a proof goal of this form (or when this is followed by other statements in sequence), use the tactic

forward_for *Inv PreIncr PostCond*

where *Inv*, *PreIncr*, *PostCond* are assertions (in PROP/LOCAL/SEP form):

Inv is the loop invariant, that holds immediately after the *init* command is executed and before each time the *test* is done; *PreIncr* is the invariant that holds immediately after the loop *body* and right before the *incr*;

PostCond is the assertion that holds after the loop is complete (whether by a break statement, or the test evaluating to false).

The following feature will appear in VST version 1.5.

Many for-loops have this special form, for (*init*; *id* < *hi*; *id*++) *body* such that the expression *hi* will evaluate to the same value every time around the loop. This upper-bound expression need not be a literal constant, it just needs to be invariant. Then you can use the tactic,

forward_for_simple_bound *n* (EX *i*:Z, PROP(\vec{P}) LOCAL(\vec{Q}) SEP(\vec{R})).

where *n* is the upper bound: a Coq value of type *Z* such that *hi* will evaluate to *n*. The loop invariant is given by the expression (EX *i*:Z, PROP(\vec{P}) LOCAL(\vec{Q}) SEP(\vec{R})), where *i* is the value (in each iteration) of the loop iteration variable *id*. This tactic generates simpler subgoals than the general forward_for tactic.

When the loop has the form, `for (id=lo; id < hi; id++) body` where *lo* is a literal constant, then the `forward_for_simple_bound` tactic will generate slightly simpler subgoals.

70 Nested Loads

This experimental feature will appear in VST release 1.5.

To handle assignment statements with nested loads, such as $x[i]=y[i]+z[i]$; the recommended method is to break it down into smaller statements compatible with separation logic: $t=y[i]$; $u=z[i]$; $x[i]=t+u$;. However, sometimes you may be proving correctness of preexisting or machine-generated C programs. Verifiable C has an **experimental** nested-load mechanism to support this.

We use an expression-evaluation relation $e \Downarrow v$ which comes in two flavors:

$\text{rel_expr} : \text{expr} \rightarrow \text{val} \rightarrow \text{rho} \rightarrow \text{mpred}.$

$\text{rel_lvalue} : \text{expr} \rightarrow \text{val} \rightarrow \text{rho} \rightarrow \text{mpred}.$

The assertion $\text{rel_expr } e \ v \ \rho$ says, “expression e evaluates to value v in environment ρ and in the current memory.” The rel_lvalue evaluates the expression as an l -value, to a pointer to the data.

Evaluation rules for rel_expr are listed here:

$\text{rel_expr_const_int} : \quad \forall (i : \text{int}) \ \tau \ (P : \text{mpred}) \ (\rho : \text{environ}),$
 $P \vdash \text{rel_expr} \ (\text{Econst_int } i \ \tau) \ (\text{Vint } i) \ \rho.$

$\text{rel_expr_const_float} : \quad \forall (f : \text{float}) \ \tau \ P \ (\rho : \text{environ}),$
 $P \vdash \text{rel_expr} \ (\text{Econst_float } f \ \tau) \ (\text{Vfloat } f) \ \rho.$

$\text{rel_expr_const_long} : \quad \forall (i : \text{int64}) \ \tau \ P \ \rho,$
 $P \vdash \text{rel_expr} \ (\text{Econst_long } i \ \tau) \ (\text{Vlong } i) \ \rho.$

$\text{rel_expr_tempvar} : \quad \forall (\text{id} : \text{ident}) \ \tau \ (v : \text{val}) \ P \ \rho,$
 $\text{Map.get } (\text{te_of } \rho) \ \text{id} = \text{Some } v \rightarrow$
 $P \vdash \text{rel_expr} \ (\text{Etempvar } \text{id } \tau) \ v \ \rho.$

$\text{rel_expr_addrof} : \quad \forall (e : \text{expr}) \ \tau \ (v : \text{val}) \ P \ \rho,$
 $P \vdash \text{rel_lvalue } e \ v \ \rho \rightarrow$
 $P \vdash \text{rel_expr} \ (\text{Eaddrof } e \ \tau) \ v \ \rho.$

$\text{rel_expr_unop} : \quad \forall P \ (e_1 : \text{expr}) \ (v_1 \ v : \text{val}) \ \tau \ op \ \rho,$
 $P \vdash \text{rel_expr } e_1 \ v_1 \ \rho \rightarrow$

$\text{Cop.sem.unary.operation } op \ v_1 \ (\text{typeof } e_1) = \text{Some } v \rightarrow$
 $P \vdash \text{rel_expr } (\text{Eunop } op \ e_1 \ \tau) \ v \ \rho.$
 $\text{rel_expr.binop: } \quad \forall (e_1 \ e_2 : \text{expr}) \ (v_1 \ v_2 \ v : \text{val}) \ \tau \ op \ P \ \rho,$
 $P \vdash \text{rel_expr } e_1 \ v_1 \ \rho \rightarrow$
 $P \vdash \text{rel_expr } e_2 \ v_2 \ \rho \rightarrow$
 $(\forall \ m : \text{Memory.Mem.mem},$
 $\text{Cop.sem.binary.operation } op \ v_1 \ e \ (\text{typeof } e_1) \ v_2 \ (\text{typeof } e_2) \ m = \text{Some } v) \rightarrow$
 $P \vdash \text{rel_expr } (\text{Ebinop } op \ e_1 \ e_2 \ \tau) \ v \ \rho.$
 $\text{rel_expr.cast: } \quad \forall (e_1 : \text{expr}) \ (v_1 \ v : \text{val}) \ \tau \ P \ \rho,$
 $P \vdash \text{rel_expr } e_1 \ v_1 \ \rho \rightarrow$
 $\text{Cop.sem.cast } v_1 \ (\text{typeof } e_1) \ \tau = \text{Some } v \rightarrow$
 $P \vdash \text{rel_expr } (\text{Ecast } e_1 \ \tau) \ v \ \rho.$
 $\text{rel_expr.lvalue: } \quad \forall (a : \text{expr}) \ (\text{sh} : \text{Share.t}) \ (v_1 \ v_2 : \text{val}) \ P \ \rho,$
 $P \vdash \text{rel_lvalue } a \ v_1 \ \rho \rightarrow$
 $P \vdash \text{mapsto } sh \ (\text{typeof } a) \ v_1 \ v_2 * \text{TT} \rightarrow$
 $v_2 <> \text{Vundef} \rightarrow$
 $P \vdash \text{rel_expr } a \ v_2 \ \rho.$
 $\text{rel_lvalue.local: } \quad \forall (\text{id} : \text{ident}) \ \tau \ (b : \text{block}) \ P \ \rho,$
 $P \vdash \text{!!}(\text{Map.get } (\text{ve_of } \rho) \ \text{id} = \text{Some } (b, \tau)) \rightarrow$
 $P \vdash \text{rel_lvalue } (\text{Evar } \text{id} \ \tau) \ (\text{Vptr } b \ \text{Int.zero}) \ \rho.$
 $\text{rel_lvalue.global: } \quad \forall (\text{id} : \text{ident}) \ \tau \ (v : \text{val}) \ P \ \rho,$
 P
 $\vdash \text{!!}(\text{Map.get } (\text{ve_of } \rho) \ \text{id} = \text{None} \wedge$
 $\text{Map.get } (\text{ge_of } \rho) \ \text{id} = \text{Some } (v, \tau)) \rightarrow$
 $P \vdash \text{rel_lvalue } (\text{Evar } \text{id} \ \tau) \ v \ \rho.$
 $\text{rel_lvalue.deref: } \quad \forall (a : \text{expr}) \ (b : \text{block}) \ (z : \text{int}) \ \tau \ P \ \rho,$
 $P \vdash \text{rel_expr } a \ (\text{Vptr } b \ z) \ \rho \rightarrow$
 $P \vdash \text{rel_lvalue } (\text{Ederef } a \ \tau) \ (\text{Vptr } b \ z) \ \rho.$
 $\text{rel_lvalue.field.struct: } \quad \forall (i \ \text{id} : \text{ident}) \ \tau \ e \ (b : \text{block}) \ (z : \text{int}) \ (\text{fList} : \text{fieldlist}) \ \text{att} \ ($
 $\text{typeof } e = \text{Tstruct } \text{id} \ \text{fList} \ \text{att} \rightarrow$
 $\text{field_offset } i \ \text{fList} = \text{Errors.OK } \delta \rightarrow$
 $P \vdash \text{rel_expr } e \ (\text{Vptr } b \ z) \ \rho \rightarrow$
 $P \vdash \text{rel_lvalue } (\text{Efield } e \ i \ \tau) \ (\text{Vptr } b \ (\text{Int.add } z \ (\text{Int.repr } \delta))) \ \rho.$

The primitive nested-load assignment rule is,

Axiom `semax_loadstore`:

$$\begin{aligned} & \forall v0\ v1\ v2\ \Delta\ e1\ e2\ sh\ P\ P', \\ & \text{writable_share } sh \rightarrow \\ & P \vdash !! (tc_val\ (typeof\ e1)\ v2) \\ & \quad \&\& \text{rel_lvalue } e1\ v1 \\ & \quad \&\& \text{rel_expr } (Ecast\ e2\ (typeof\ e1))\ v2 \\ & \quad \&\& (\text{mapsto } sh\ (typeof\ e1)\ v1\ v0) * P' \rightarrow \\ & \text{semax } \Delta\ (\triangleright P)\ (\text{Sassign } e1\ e2) \\ & \quad (\text{normal_ret_assert } (\text{mapsto } sh\ (typeof\ e1)\ v1\ v2) * P')). \end{aligned}$$

but do not use this rule! It is best to use a derived rule, such as,

Lemma `semax_loadstore_array`:

$$\begin{aligned} & \forall n\ vi\ lo\ hi\ t1\ (\text{contents: } Z \rightarrow \text{reptype } t1)\ v1\ v2\ \Delta\ e1\ ei\ e2\ sh\ P\ Q\ R, \\ & \text{reptype } t1 = \text{val} \rightarrow \\ & \text{type_is_by_value } t1 \rightarrow \\ & \text{legal_alignas_type } t1 = \text{true} \rightarrow \\ & \text{typeof } e1 = \text{tptr } t1 \rightarrow \\ & \text{typeof } ei = \text{tint} \rightarrow \\ & \text{PROP}_x P\ (\text{LOCAL}_x Q\ (\text{SEP}_x R)) \\ & \quad \vdash \text{rel_expr } e1\ v1 \\ & \quad \&\& \text{rel_expr } ei\ (\text{Vint } (\text{Int.repr } vi)) \\ & \quad \&\& \text{rel_expr } (Ecast\ e2\ t1)\ v2 \rightarrow \\ & \text{nth_error } R\ n = \text{Some } (\text{array_at } t1\ sh\ \text{contents } lo\ hi\ v1)) \rightarrow \\ & \text{writable_share } sh \rightarrow \\ & \text{tc_val } t1\ v2 \rightarrow \\ & \text{in_range } lo\ hi\ vi \rightarrow \\ & \text{semax } \Delta\ (\triangleright \text{PROP}_x P\ (\text{LOCAL}_x Q\ (\text{SEP}_x R))) \\ & \quad (\text{Sassign } (\text{Ederef } (\text{Ebinop } \text{Oadd } e1\ ei\ (\text{tptr } t1))\ t1)\ e2) \\ & \quad (\text{normal_ret_assert} \\ & \quad (\text{PROP}_x P\ (\text{LOCAL}_x Q\ (\text{SEP}_x \\ & \quad (\text{replace_nth } n\ R \\ & \quad \quad \text{array_at } t1\ sh\ (\text{upd } \text{contents } vi\ (\text{valinject } _ v2))\ lo\ hi\ v1)))))). \end{aligned}$$

Proof-automation support is available for `semax_loadstore_array` and `rel_expr`, in the form of the `forward_n1` (for “forward nested loads”) tactic. For example, with this proof goal,

```
semax Delta
(PROP ()
  LOCAL(`(eq (Vint (Int.repr i))) (eval_id _i); `(eq x) (eval_id _x);
    `(eq y) (eval_id _y); `(eq z) (eval_id _z))
  SEP(`(array_at tdouble Tsh (Vfloat oo fx) 0 n x);
    `(array_at tdouble Tsh (Vfloat oo fy) 0 n y);
    `(array_at tdouble Tsh (Vfloat oo fz) 0 n z)))
(Ssequence
  (Sassign (* x[i] = y[i] + z[i]; *)
    (Ederef (Ebinop Oadd (Etempvar _x (tptr tdouble)) (Etempvar _i tint)
      (tptr tdouble)) tdouble)
    (Ebinop Oadd
      (Ederef (Ebinop Oadd (Etempvar _y (tptr tdouble)) (Etempvar _i tint)
        (tptr tdouble)) tdouble)
      (Ederef (Ebinop Oadd (Etempvar _z (tptr tdouble)) (Etempvar _i tint)
        (tptr tdouble)) tdouble) tdouble))
    MORE_COMMANDS)
  POSTCONDITION
```

the tactic-application `forward_n1` yields the new proof goal,

```
semax Delta
(PROP ()
  LOCAL(`(eq (Vint (Int.repr i))) (eval_id _i); `(eq x) (eval_id _x);
    `(eq y) (eval_id _y); `(eq z) (eval_id _z))
  SEP
    (`(array_at tdouble Tsh
      (upd (Vfloat oo fx) i (Vfloat (Float.add (fy i) (fz i)))) 0 n x);
      `(array_at tdouble Tsh (Vfloat oo fy) 0 n y);
      `(array_at tdouble Tsh (Vfloat oo fz) 0 n z)))
  MORE_COMMANDS
  POSTCONDITION
```

71 JUNK

$$\begin{array}{c}
\text{makelock} \frac{R \text{ positive} \quad R \text{ precise}}{\Delta \vdash \{l \mapsto 0\} \text{makelock } l \{l \sqsupset^\bullet R\}} \\
\\
\text{freelock} \frac{}{\Delta \vdash \{R * l \sqsupset^\bullet R\} \text{freelock } l \{R * l \mapsto 0\}} \\
\\
\text{splitlock} \frac{\pi_1 \oplus \pi_2 = \pi}{l \sqsupset^{\pi_1} R * l \sqsupset^{\pi_2} R \leftrightarrow l \sqsupset^\pi R} \\
\\
\text{acq} \frac{}{\Delta \vdash \{l \sqsupset^\pi R\} \text{acquire } l \{R * l \sqsupset^\pi R\}} \\
\\
\text{rel} \frac{}{\Delta \vdash \{R * l \sqsupset^\pi R\} \text{release } l \{l \sqsupset^\pi R\}} \\
\\
\text{spawn} \frac{}{\Delta \vdash \{(a : \{\vec{y}P\}\{\text{emp}\}) \wedge (Px)[\vec{b}/\vec{y}] * F\} \text{spawn } a(\vec{b}) \{F\}} \\
\\
\text{exit} \frac{}{\Delta \vdash \{\text{emp}\} \text{exit}() \{\perp\}}
\end{array}$$

Figure 71.1: CSL rules for threads and locks